Sequential Half-Aggregation of Lattice-Based Signatures*

Katharina Boudgoust¹ and Akira Takahashi²

¹ Aarhus University, Denmark katharina.boudgoust@cs.au.dk ² University of Edinburgh, UK takahashi.akira.58s@gmail.com

 $March\ 4,\ 2024$

Abstract. With Dilithium and Falcon, NIST selected two lattice-based signature schemes during their post-quantum standardization project. Whereas Dilithium follows the Fiat-Shamir with Aborts (Lyubashevsky, Asiacrypt'09) blueprint, Falcon can be seen as an optimized version of the GPV-paradigm (Gentry et al., STOC'06). An important question now is whether those signatures allow additional features such as the aggregation of distinct signatures. One example are sequential aggregate signature (SAS) schemes (Boneh et al., Eurocrypt'04) which allow a group of signers to sequentially combine signatures on distinct messages in a compressed manner. The present work first reviews the state of the art of (sequentially) aggregating lattice-based signatures, points out the insecurity of one of the existing Falconbased SAS (Wang and Wu, PROVSEC'19), and proposes a fix for it. We then construct the first Fiat-Shamir with Aborts based SAS by generalizing existing techniques from the discrete-log setting (Chen and Zhao, ESORICS'22) to the lattice framework. Going from the pre-quantum to the post-quantum world, however, does most often come with efficiency penalties. In our work, we also meet obstacles that seem inherent to lattice-based signatures, making the resulting scheme less efficient than what one would hope for. As a result, we only achieve quite small compression rates. We compare our construction with existing lattice-based SAS which all follow the GPV-paradigm. The bottom line is that none of the schemes achieves a good compression rate so far.

^{*} An extended abstract appeared at ESORICS 2023. This is the full version.

Table of Contents

1	Introduction			
	1.1 Our Contributions	4		
	1.2 Other Related Work	5		
2	Preliminaries	5		
	2.1 Probability and Regularity	5		
	2.2 Module Lattice Problems	6		
	2.3 Fiat-Shamir with Aborts Signatures	6		
	2.4 Sequential Aggregate Signatures	7		
3	Sequential Half-Aggregation of FSwA Signatures			
	3.1 Definition and Correctness of the Scheme	8		
	3.2 Security Proof	10		
4	Performance Estimates and Comparison			
	4.1 Performance Estimates	15		
	4.2 Comparison With SAS Using Trapdoors	15		
5	Attacks on Existing Schemes	17		
	5.1 Attack on [WW19]	17		
	5.2 Attack on [FH20]	19		
6	Conclusion	21		
A	More Security Notions for Sequential Aggregate Signatures	24		
	A.1 Histrory-Free Sequential Aggregate Signatures	24		
	A.2 Partial-Signature History-Free Sequential Aggregate Signatures	24		

1 Introduction

Aggregate signature (AS) schemes, introduced by [BGLS03], allow N signers to individually produce signatures $\sigma_1, \ldots, \sigma_N$ on distinct messages m_1, \ldots, m_N , and later combine them into a single, compact signature σ_{AS} . Such σ_{AS} can be verified with respect to the participants' verification keys $\mathsf{pk}_1, \ldots, \mathsf{pk}_N$. Classical applications of aggregate signatures include certificate chains: in a public key infrastructure one has to include their certificate in every sent message, which itself comes from a chain of certificates issued by different authorities. Since the naive concatenation of single-user signatures significantly adds to the certificate chain (e.g., [BGLS03] reports 15% of a typical X.509 certificate length is occupied by the signature), it is paramount to replace them with a compact, aggregated signature to save bandwidth. In the literature, essentially two different paradigms of fully compact aggregate signatures have been proposed: (1) dedicated constructions based on bilinear pairings [BGLS03, BNN07], and (2) generic solutions exploiting iO [HKW15] or non-interactive arguments [DGKV22, WW22, ACL⁺22], where a signature aggregator produces a succinct proof of knowledge of N valid signatures.

There also exists the slightly restricted notion of sequential aggregate signatures (SAS) [LMRS04]. In this setting, signing and aggregation are carried out altogether: signer i associated with pk_i receives from signer i-1 aggregate so-far σ_{i-1} with a key-message list $L_{i-1} = (\mathsf{pk}_1, m_1, \ldots, \mathsf{pk}_{i-1}, m_{i-1})$, adds a signature on the message m_i of their own choice to produce σ_i , and then passes along σ_i and $L_i = L_{i-1}||(\mathsf{pk}_i, m_i)|$ to the next signer i+1. Unlike general aggregate signatures, SAS require round-robin communication among signers, which however fits well in typical application scenarios such as a certificate chain. A plethora of work proposed highly efficient, constant-size SAS using pairings [LOS+06, BGOY07, BNN07, FLS12] or assuming the existence of trapdoor permutations [LMRS04, Nev08, BGR12, GOR18].

Half-Aggregation of Fiat-Shamir Signatures. Perhaps unsurprisingly, not many aggregation methods tailored to Fiat-Shamir signatures [FS87] such as Schnorr [Sch91] are known.³ Fiat-Shamir signatures are typically constructed from three-round Σ -protocols [Cra96]: the signer invokes the underlying Σ -protocol prover to generate the first-round commit value u, samples random challenge c by hashing u together with the message m to be signed, creates response z, and outputs $\sigma = (c, z)$ as a signature. The verifier then reconstructs u from (pk, c, z) through certain algebraic operations and checks the recomputed hash against c. Equivalently, the signer can set $\sigma = (u, z)$ and the verifier recomputes the hash c, while checking if a certain relation between c and (pk, u, z) holds. The difficulty of aggregating Fiat-Shamir mainly lies in the challenge hash function: since its typical instantiation such as SHA-256 has no algebraic structure, it does not blend well with nice homomorphic properties of the underlying Σ -protocol transcript. This is why the existing approaches (e.g., [BN06, DEF+19, NRS21]) require (at least) two rounds of interaction so that all signers can first agree on a combined u that leads to the same challenge c, from which they compute shares of z.

To avoid interaction, recent papers proposed half-aggregation of Schnorr/EdDSA [CGKN21, Kas22, CZ22]. These are middle ground solutions where only the u or the z component gets aggregated, and the other part consists of a concatenation of N partial signatures. Although it is asymptotically no better than the trivial concatenation of N signatures, reducing the signature size by a constant factor has meaningful implications in practice, e.g., in certain cases the entire certificate chain of size O(N) needs to be transmitted anyway.

Another possible approach would be adapting one of the aforementioned generic solutions and having an aggregator node to prove the knowledge of N tuples of the form (u,c,z) satisfying the verification conditions descried as a circuit. However, the prover's complexity likely hinges on mixture of algebraic operations and non-algebraic hash computation in verification. Another issue with such a generic solution applied to typical Fiat-Shamir signatures is that the security proof would likely rely on heuristics. Since an aggregator of the generic method requires a concrete description of the hash function, its security is only guaranteed assuming the security of the underlying signature

³ It is well known that *interactive* multi-signatures can be generically converted to interactive aggregate signatures by asking all participants to sign a concatenation of N messages and public keys [BN06, DEF⁺19, NRS21]. However, this requires the signers to agree on all N messages and who they co-sign with in advance, and does not fit in the typical use cases of aggregate signatures such as a certificate chain.

scheme in the standard model, whereas the majority of existing Fiat-Shamir signatures are only proven in the random oracle model.

Aggregate Signatures from Lattices. Given that NIST has announced in their post-quantum cryptography standardization project two signature finalists, Falcon [PFH⁺20] and Dilithium [LDK⁺20], based on (structured) lattice assumptions, a natural question is whether tailor-made aggregate signatures can be instantiated using lattices (instead of generic solutions such as [DGKV22, ACL⁺22]). Both finalists represent the two major design principles to build lattice-based signatures: Dilithium follows Lyubashevsky's Fiat-Shamir with Aborts (FSwA) paradigm [Lyu09, Lyu12] and Falcon is a GPV-type signature using preimage sampleable trapdoor functions [GPV08].

There are a limited number of proposals within the FSwA paradigm. Boneh and Kim [BK20] presented a lattice-based instantiation of [BN06] but it requires three rounds of interactions. Boudgoust and Roux-Langlois [BR21] are the first to securely instantiate non-interactive half-aggregation of FSwA. From a high level perspective, they adapt the half-aggregation of Schnorr [CGKN21] to the lattice-setting. Whereas in Schnorr, it does not really matter whether we output $\sigma = (u, z)$ or $\sigma = (c, z)$, it makes a big difference in the lattice setting. The signature size significantly decreases in the second case. During the half-aggregation of [BR21], only the z-parts are aggregated, but all the u-parts are transmitted. Note that it is not sufficient to transmit all the c-parts, as we cannot recover the different commitments anymore from an aggregated response. However, we need every single commitment in order to verify an aggregate signature. In consequence, the provably secure version of their construction outputs a signature $\sigma_{AS} = (u_1, \dots, u_N, z)$ which is always larger than the naive concatenation of N signatures $\sigma_{con} = (c_1, z_1, \dots, c_N, z_N)$. The MMSAT scheme [DHSS20] is a candidate half-aggregate signature scheme based on a non-standard lattice problem, called the Partial Fourier Recovery problem. However, it turned out that the security proof is flawed and even simple forgery attacks exist [BR21]. Regarding sequential aggregation, the only known lattice-based solutions we are aware of follow the GPV-paradigm [EB14, WW19], of which the latter turns out to be insecure as we sketch below. Given all this, we are motivated to ask the following question in this paper:

Can we construct a non-interactive sequential half-aggregate FSwA signature scheme (1) with a signature size smaller than the naive concatenation, and (2) without invoking expensive generic solutions?

1.1 Our Contributions

In this work, we positively answer this question. In Section 3, we present a sequential half-aggregate signature based on the Fiat-Shamir with Aborts framework. The aggregation paradigm closely follows recent Schnorr-based SAS due to Chen and Zhao [CZ22]. As elaborated before, the main obstacle in previous works is that without interaction it is difficult to aggregate the commitments u_1, \ldots, u_N that are responsible for the large aggregate signature size in [BR21]. If, however, we place ourselves in the sequential aggregate model, we can aggregate over the u-parts by letting the parties sign one after each other. A sequential aggregate signature of our construction now is of the form $\sigma_{SAS} = (u, z_1, \ldots, z_N)$. Once the size of (c_1, \ldots, c_N) is larger than the size of u, our SAS produces signatures that are indeed smaller than the trivial concatenation σ_{con} . Unfortunately, when looking at the ratio between σ_{SAS} and σ_{con} , it is the (z_1, \ldots, z_N) -part (that both have in common) that makes up for most of the signature size and hence the compression rate is close to 1. Although our concrete parameter estimates in Section 4 indicate the output signature is only $\sim 1\%$ smaller than the naive concatenation, we believe ours to be an important step towards better understanding the possibilities and limits of lattice-based aggregate signatures.

The security of our scheme tightly reduces to the existential unforgeability of the standard single-user FSwA scheme instantiated with structured lattices. We prove security in the so-called full history setting of SAS. For completeness, we also discuss its security in a new model that has been introduced in [CZ22], which we call the partial-signature history-free security model. The details can be found in Appendix A. Although our construction closely follows the one of [CZ22], our security proof is more involved because of subtleties that arise in the lattice setting. We have to consider several bad events that might happen and bound their probability. In Section 4, we also compare our scheme with the two existing lattice-based SAS [EB14, WW19] following the

GPV-paradigm. As in the lattice setting we only have so-called preimage sampleable trapdoor functions (and no trapdoor permutation), they cannot achieve constant-size SAS either. The upshot is that neither of them saves more than 4% of signature size if a fair comparison is made against the naive concatenation and taking recent advances [ETWY22] into account.

As a separate contribution, we point out insecurities of two existing aggregate signature schemes explicitly instantiated with NIST finalists: (1) Falcon-based SAS of [WW19] does not guarantee the claimed security property due to the existence of a forgery attack (Section 5.1), and (2) Dilithium-based interactive multi-signature of [FH20] (which can be generically turned into an aggregate signature using [BN06]'s trick) leaks part of the secret key due to the misuse of Bai-Galbraith's HighBits optimization trick [BG14] (Section 5.2). The latter attack highlights that, even after knowing how to aggregate signatures that follow a general paradigm, it is not trivial to make the aggregation work for optimized instantiations.

Given the above attacks on existing solutions and the concrete parameter estimates of our construction, we conclude in Section 6 that concretely efficient aggregation of the lattice-based NIST finalists is still an unexplored area and mark it as an interesting direction for future work.

1.2 Other Related Work

Imposing a sequential way of signing is not the only way how to restrict the model of aggregate signatures. Other works look for instance at a synchronous model [GR06, AGH10], where signatures are aggregated together if they have been issued at the same time interval. One recent result studies lattice-based aggregate signatures in the synchronous model [FSZ22]. As already mentioned before, a related concept are multi-signatures [MOR01, Bol03], where we allow the parties to interact with each other. There are several recent results on lattice-based multi-signatures [DOTT21, BTT22] and we refer to the references therein.

2 Preliminaries

Notations. For any positive integer N, we denote by [N] the set $\{1,\ldots,N\}$. For a finite set S, we denote its cardinality by [S] and the uniform distribution over S by U(S). We simply write $s \overset{\$}{\leftarrow} S$ to indicate sampling s from U(S). For a probability distribution \mathscr{D} , we write $s \overset{\$}{\leftarrow} \mathscr{D}$ to indicate sampling s from \mathscr{D} ; for a randomized (resp. deterministic) algorithm \mathscr{A} we write $s \overset{\$}{\leftarrow} \mathscr{A}$ (resp. $s := \mathscr{A}$) to indicate assigning an output from \mathscr{A} to s. Throughout, the security parameter is denoted by λ . The abbreviation PPT stands for probabilistic polynomial-time.

Following [LDK+20] $r' = r \mod^+ \alpha$ denotes the unique integer $r' \in [0, \alpha)$ such that $r' \equiv r \mod \alpha$. For an even (resp. odd) positive integer α , $r' = r \mod^{\pm} \alpha$ denotes the unique integer $r' \in (-\alpha/2, \alpha/2]$ (resp. $r' \in [-(\alpha-1)/2, (\alpha-1)/2]$) such that $r' \equiv r \mod \alpha$.

Throughout the paper, we work over the ring $R =: \mathbb{Z}[X]/\langle X^n + 1 \rangle$, where n is a power of 2. For any ring element $r \in R$, we define $||r||_2$, $||r||_1$ and $||r||_{\infty}$ to be the respective norms of its coefficient vector. For some prime q, we define $R_q := R/(qR)$ and for some positive integer γ , we set $S_{\gamma} := \{r \in R : ||r||_{\infty} \leq \gamma\}$.

2.1 Probability and Regularity

The Rényi divergence (RD) defines a measure of distribution closeness. We follow [BLR⁺18] and set the RD as the exponential of the classical definition.

Definition 2.1 (Rényi Divergence). Let P and Q be two discrete probability distributions such that $Supp(P) \subseteq Supp(Q)$. The Rényi divergence (of order 2) is defined by

$$\mathrm{RD}_2(P\|Q) = \sum_{x \in \mathsf{Supp}(P)} \frac{P(x)^2}{Q(x)}.$$

The RD fulfills the following properties, as proved in [vEH14].

Lemma 2.2. Let P,Q be two discrete probability distributions with $\mathsf{Supp}(P) \subseteq \mathsf{Supp}(Q)$.

Data Processing Inequality: $RD_2(P^g||Q^g) \leq RD_2(P||Q)$ for any function g, where P^g (resp. Q^g) denotes the distribution of g(y) induced by sampling $y \leftarrow P$ (resp. $y \leftarrow Q$),

Probability Preservation: Let $E \subset \mathsf{Supp}(Q)$ be an event, then

$$P(E) \le \sqrt{Q(E) \cdot \text{RD}_2(P||Q)}$$

In Section 3, we need the following regularity result.

Lemma 2.3 ([BJRW23, Lem. 2.8] Simplified). Let $k, \ell, q, \gamma \in \mathbb{N}$ such that q is prime. Further, let $R = \mathbb{Z}[X]/\langle X^n + 1 \rangle$, where n is a power of 2. Then,

$$\mathrm{RD}_2((\mathbf{A}, \mathbf{A}\mathbf{y}') \| (\mathbf{A}, \mathbf{v})) \le \left(1 + \frac{q^k}{(2\gamma + 1)^\ell}\right)^n,$$

where $\mathbf{A} \stackrel{\$}{\leftarrow} R_q^{k \times \ell}$, $\mathbf{y}' \stackrel{\$}{\leftarrow} \mathscr{D} := U(S_{\gamma}^{\ell})$ and $\mathbf{v} \stackrel{\$}{\leftarrow} R_q^k$.

In order to obtain a constant Rényi divergence, we require

$$\ell \ge k \cdot \frac{\log_2 q}{\log_2(2\gamma + 1)} + O\left(\frac{\log_2 q}{\log_2(2\gamma + 1)}\right). \tag{1}$$

Remark 2.4. Alternatively, if one prefers the discrete Gaussian distribution for \mathscr{D} , one can use the regularity result [LPR13, Cor. 7.5]. It comes with the advantage that it holds for any $q \geq 2$, not necessarily prime, and that the parameter ℓ can be arbitrarily close to k, in particular $\ell = k$ is a possible choice, which is common in practice. However, the resulting Gaussian width, defining the parameter γ , cannot be freely chosen and needs to be above a certain threshold, which is much larger than parameters used in practice.

2.2 Module Lattice Problems

We also recall two lattice problems and refer to [LS15] for more details. We state them in their respective discrete, primal and HNF form.

Definition 2.5 (M-LWE). Let $k, \ell, \eta \in \mathbb{N}$. The Module Learning With Errors problem M-LWE_{k,ℓ,η} is defined as follows. Given $\mathbf{A} \overset{\$}{\leftarrow} R_q^{k \times \ell}$ and $\mathbf{t} \in R_q^k$. Decide whether $\mathbf{t} \overset{\$}{\leftarrow} R_q^k$ or if $\mathbf{t} = [\mathbf{A} | \mathbf{I}_k] \cdot \mathbf{s}$, where $\mathbf{s} \overset{\$}{\leftarrow} S_q^{\ell+k}$.

The M-LWE assumption states that no PPT algorithm can distinguish between the two distributions with non-negligible advantage. Worst-case to average-case reductions guarantee that M-LWE is quantumly [LS15] and classically [BJRW20] at least as hard as the approximate shortest vector problem over module lattices.

Definition 2.6 (M-SIS). Let $k, \ell, b \in \mathbb{N}$. The Module Short Integer Solution problem M-SIS_{k,\ell,b} is as follows. Given a uniformly random matrix $\mathbf{A} \stackrel{\$}{\leftarrow} R_q^{k \times \ell}$. Find a non-zero vector $\mathbf{s} \in R_q^{k+\ell}$ such that $\|\mathbf{s}\|_2 \leq b$ and $[\mathbf{A}|\mathbf{I}_k] \cdot \mathbf{s} = \mathbf{0} \in R_q^k$.

The M-SIS assumption states that no PPT adversary can solve this problem with non-negligible probability. Worst-case to average-case reductions guarantee that M-SIS is classically [LS15] at least as hard as the approximate shortest independent vector problem over module lattices.

2.3 Fiat-Shamir with Aborts Signatures

In this paper, we build a sequential aggregate signature FSwA-SAS starting from a well-studied signature scheme FSwA-S = (Setup, Gen, Sign, Ver) whose definition we recall in Algorithm 1. It follows the so-called *Fiat-Shamir with Aborts* paradigm [Lyu09, Lyu12] and can be seen as the module variant of [GLP12] or the 'vanilla' flavor of Dilithium.

```
Algorithm 1: Description of the FSwA-S Signature
The challenge space is \mathsf{Ch} := \{c \in \mathbb{R} : ||c||_{\infty} = 1 \land ||c||_{1} = \kappa \} and the message space is M = \{0,1\}^{\ell}.
    The random oracle is H: \{0,1\}^* \to Ch.
\mathsf{Setup}(1^{\lambda})
                                                                                                                 Gen(\bar{\mathbf{A}})
                                                                                                                         1: \mathbf{s} \stackrel{\$}{\leftarrow} S_n^{\ell+k}
       1: \mathbf{A} \overset{\$}{\leftarrow} R_a^{k \times \ell}
       \mathbf{\bar{A}} := [\mathbf{\bar{A}}|\mathbf{I}_k]
                                                                                                                         2: \mathbf{t} := \bar{\mathbf{A}}\mathbf{s} \bmod q
                                                                                                                         3: \mathsf{sk} := \mathbf{s}
       3: return Ā
                                                                                                                         4: \mathsf{pk} := \mathbf{t}
\mathsf{Sign}(\mathsf{sk}, m)
                                                                                                                         5: return (sk, pk)
       1: \mathbf{s} := \mathsf{sk}
                                                                                                                 Ver(pk, \sigma, m)
        2: \mathbf{t} := \bar{\mathbf{A}}\mathbf{s} \mod q
                                                                                                                         1: (\mathbf{u}, \mathbf{z}) := \sigma
       3: \mathbf{z} := \bot
        4: while \mathbf{z} := \bot \mathbf{do}
                                                                                                                         2: \mathbf{t} := \mathsf{pk}
                   \mathbf{y} \leftarrow \mathscr{D}^{\ell+k}
                                                                                                                         3: c := \mathsf{H}(\mathbf{u}, \mathbf{t}, m)
                   \mathbf{u} := \bar{\mathbf{A}}\mathbf{y} \bmod q
                                                                                                                         4: \|\mathbf{z}\|_{\infty} \leq B \wedge \bar{\mathbf{A}}\mathbf{z} = c \cdot \mathbf{t} + \mathbf{u} \text{ then}
                   c := \mathsf{H}(\mathbf{u}, \mathbf{t}, m)
                                                                                                                         5: return 1
                                                                                                                         6: else
                   \mathbf{z} := c \cdot \mathbf{s} + \mathbf{y}
                   z := RejSamp(z, c \cdot s)
                                                                                                                                     return 0
      10: \sigma := (\mathbf{u}, \mathbf{z})
     11: return \sigma
```

Modification. A difference to the standard design is that instead of outputting $\sigma = (c, \mathbf{z})$, we output $\sigma = (\mathbf{u}, \mathbf{z})$. For a single signature, both cases are equivalent, as \mathbf{u} defines c via the hash function H (and the public key \mathbf{t} and the message m) and c defines \mathbf{u} via the equation $\mathbf{u} = \bar{\mathbf{A}}\mathbf{z} - c \cdot \mathbf{t}$. However, this is not the case for a (sequential) aggregate signature scheme and we thus need to transmit the information \mathbf{u} .

Distribution \mathscr{D} . During the signing algorithm Sign, the FSwA-S scheme uses a distribution \mathscr{D} to sample a vector of ring elements of short norm over R. In the literature, mainly two different ways of instantiating \mathscr{D} are studied. The first uses discrete Gaussian distributions (as for instance in [Lyu12]) and the second uses the uniform distribution over a bounded set, i.e., $\mathscr{D} = U(S_{\gamma})$ for some $\gamma \ll q$ (as for instance in [GLP12]). The concrete instantiation of \mathscr{D} then influences the choice of the rejection algorithm RejSamp during signing and of the bound B during verification. In this paper, we focus on the latter as this is the choice commonly used in practice, as for instance in Dilithium. In this case, the algorithm RejSamp outputs \bot if $\|\mathbf{z}\|_{\infty} > \gamma - \kappa \cdot \eta =: B$, else it outputs \mathbf{z} .

Security. Overall, the UF-CMA security of the scheme FSwA-S as specified in Algorithm 1 is based on the hardness of M-LWE and M-SIS [LS15]. For the reason of space limits, we refer the interested reader to the original security proofs in [Lyu12, GLP12] in the random oracle model.

2.4 Sequential Aggregate Signatures

Sequential aggregate signatures (SAS) were first introduced in [LMRS04]. We recall now the syntax of a (full-history) SAS scheme, together with the definitions of correctness and security following the notations of Gentry et al. [GOR18]. See also Appendix A for a discussion of alternative definitions of SAS schemes and corresponding security notions.

Definition 2.7 (SAS). A sequential aggregate signature scheme (SAS) for a message space M consists of a tuple of PPT algorithms SAS = (Setup, Gen, SeqSign, SeqVerify) defined as follows:

Setup(1^{λ}) \rightarrow pp: On input the security parameter λ , the setup algorithm outputs the public parameters pp.

Gen(pp) → (sk, pk): On input the public parameters pp, the key generation algorithm outputs a pair of secret key sk and public key pk.

```
Game 1: Description of the FH-UF-CMA<sub>SAS</sub>(\mathcal{A}, \lambda) Security Game
1: pp \leftarrow Setup(1^{\lambda})
                                                                                            \mathsf{OSeqSign}(m_i, \sigma_{i-1}, L_{i-1})
2: (pk, sk) \leftarrow Gen(pp)
                                                                                                   1: \sigma_i \leftarrow \mathsf{SeqSign}(\mathsf{sk}, m_i, L_{i-1}, \sigma_{i-1})
3: \mathcal{Q} := \emptyset
                                                                                                  2: Q := Q \cup \{(m_i, L_{i-1})\}
4: (L_N^*, \sigma_N^*) \leftarrow \mathcal{A}^{\mathsf{OSeqSign}}(\mathsf{pp}, \mathsf{pk})
                                                                                                  3: return \sigma_i
5: if SeqVerify (L_N^*, \sigma_N^*) \wedge \exists i^* \in [N]: (\mathsf{pk}_{i^*} =
     pk \wedge (m_{i^*}, L_{i^*-1}) \notin \mathcal{Q}) then
6:
          {\bf return}\ 1
7: else
           return 0
8:
```

SeqSign($\operatorname{sk}_i, m_i, L_{i-1}, \sigma_{i-1}$) $\rightarrow \sigma_i$: On input a secret key sk_i , a message $m_i \in M$, a list L_{i-1} with $L_{i-1} := (\operatorname{pk}_1, m_1) || \ldots || (\operatorname{pk}_{i-1}, m_{i-1})$, and a so-far signature σ_{i-1} , the sequential signing algorithm outputs a new so-far signature σ_i .

SeqVerify $(L_N, \sigma_N) \to \{0, 1\}$: On input a list L_N of N message-public-key pairs and a sequential aggregate signature σ_N , the sequential verification algorithm either outputs 1 (accept) or 0 (reject).

For convenience, given a list $L_j = (\mathsf{pk}_1, m_1) || \dots || (\mathsf{pk}_j, m_j)$, we denote by L_i its *i*th prefix $L_i := (\mathsf{pk}_1, m_1) || \dots || (\mathsf{pk}_i, m_i)$ for $1 \le i < j$.

Definition 2.8 (Correctness). Let SAS = (Setup, Gen, SeqSign, SeqVerify) be a sequential aggregate signature scheme for a message space M. It is called correct if for all $\lambda, N \in \mathbb{N}$ it yields

$$\Pr\left[\mathsf{SeqVerify}(L_N, \sigma_N) = 1\right] = 1 - \mathsf{negl}(\lambda),$$

 $\label{eq:where mi} \begin{aligned} \textit{where } m_i \in M, \ \mathsf{pp} \leftarrow \mathsf{Setup}(1^\lambda), \ (\mathsf{sk}_i, \mathsf{pk}_i) \leftarrow \mathsf{Gen}(\mathsf{pp}), \ L_i = (\mathsf{pk}_1, m_1) || \ldots || (\mathsf{pk}_i, m_i) \ \textit{ and } \ \sigma_i \leftarrow \mathsf{SeqSign}(\mathsf{sk}_i, m_i, L_{i-1}, \sigma_{i-1}) \ \textit{ for all } i \in [N]. \ \textit{Let } L_0 = \varnothing \ \textit{ and } \ \sigma_0 = (\mathbf{0}, \mathbf{0}). \end{aligned}$

Informally, full history unforgeability against chosen message attacks captures the following security notion. An adversary is given a challenge public key and has access to a sequential signing oracle that, on input a message, a so-far signature and a list of public keys and messages (called 'history') provides the next so-far signature using the secret key corresponding to the challenge key. A forgery is composed of an sequentially aggregate signature together with a history (i.e., a list of message-key pairs). The forgery is successful if it passes verification, if one of the public keys is the challenge key and if the signing oracle has not yet been queried on the same message and history.

Definition 2.9 (FH-UF-CMA Security). A SAS scheme satisfies <u>full history unforgeability against chosen message attacks</u>, if for all PPT adversaries A,

$$\mathsf{Adv}_{\mathsf{SAS}}^{\mathsf{FH}\mathsf{-}\mathsf{UF}\mathsf{-}\mathsf{CMA}}(\mathcal{A}) := \Pr\left[\mathsf{FH}\mathsf{-}\mathsf{UF}\mathsf{-}\mathsf{CMA}_{\mathsf{SAS}}(\mathcal{A},\lambda) = 1\right] = \mathsf{negl}(\lambda),$$

where the FH-UF-CMA_{SAS} game is described in Game 1.

3 Sequential Half-Aggregation of FSwA Signatures

3.1 Definition and Correctness of the Scheme

Our scheme is described in Algorithm 2. The overall structure closely follows the one by Chen and Zhao [CZ22]. We remark that, for the sake of security proof, the key generation algorithm slightly differs from the original one in Algorithm 1. It keeps regenerating a key pair until the public key \mathbf{t} contains at least one invertible element. This terminates relatively quickly in practice. Let p_{inv} be the probability that $\mathbf{t} = \bar{\mathbf{A}}\mathbf{s}$ has at least one invertible coefficient over R_q , where \mathbf{s} is uniformly sampled from $S_{\eta}^{\ell+k}$. Then the expected running time of Gen is $1/p_{\mathsf{inv}}$. One can experimentally find p_{inv} for each parameter set.

The single signature scheme FSwA-S as described in Algorithm 1 can be obtained from the sequential aggregate signature scheme FSwA-SAS (Algorithm 2) when taking N=1 and starting with the empty set $L_0=\varnothing$ and a trivial so-far signature $\sigma_0=(\mathbf{0},\mathbf{0})$.

```
Algorithm 2: Description of the FSwA-SAS Sequential Aggregate Signature
The challenge space is \mathsf{Ch} := \{c \in R : \|c\|_{\infty} = 1 \land \|c\|_{1} = \kappa\} and the message space is M' = \{0,1\}^{l}. The random oracle is \mathsf{H} : \{0,1\}^{*} \to \mathsf{Ch}. The starting point is i=1. Let L_0 = \varnothing and \sigma_0 = (\mathbf{0},\mathbf{0}).
     Setup is as in Algorithm 1.
  Gen(\bar{\mathbf{A}})
                                                                                                            SeqVerify(L_N, \sigma_N)
       1: t := 0
                                                                                                                   1: (\mathbf{t}_1, m_1) || \dots || (\mathbf{t}_N, m_N) := L_N
       2: while t has no invertible coefficient do
                                                                                                                   2: (\tilde{\mathbf{u}}_N, \mathbf{z}_1, \dots, \mathbf{z}_N) := \sigma_N
                     \mathbf{s} \stackrel{\$}{\leftarrow} S_n^{\ell+k}
                                                                                                                   3: \mathbf{z}_0 := 1
                     \mathbf{t} := \bar{\mathbf{A}}\mathbf{s} \bmod q
                                                                                                                   4: if \exists i such that \mathbf{t}_i has no invertible element
                                                                                                                         then
       5: sk := s
                                                                                                                           return 0
       6: pk := t
                                                                                                                   6: for i = N, ..., 1 do
       7: return (sk, pk)
                                                                                                                               if \|\mathbf{z}_i\|_2 > B then
 \mathsf{SeqSign}(\mathsf{sk}_i, m_i, L_{i-1}, \sigma_{i-1})
                                                                                                                   8:
                                                                                                                                       return 0
       1: (\tilde{\mathbf{u}}_{i-1}, \mathbf{z}_1, \dots, \mathbf{z}_{i-1}) := \overline{\sigma_{i-1}}
                                                                                                                  9:
                                                                                                                                  L_i := (\mathbf{t}_1, m_1) || \dots || (\mathbf{t}_i, m_i) ||
       \mathbf{s}_i := \mathsf{sk}_i
                                                                                                                 10:
                                                                                                                                  \overline{c_i} := \mathsf{H}(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})
       3: \mathbf{t}_i := \bar{\mathbf{A}}\mathbf{s}_i \bmod q
                                                                                                                                 \mathbf{u}_i := \bar{\mathbf{A}} \mathbf{z}_i - c_i \mathbf{t}_i \mod q
\tilde{\mathbf{u}}_{i-1} := \tilde{\mathbf{u}}_i - \mathbf{u}_i \mod q
                                                                                                                 11:
       4: L_i := L_{i-1} || (\mathbf{t}_i, m_i) |
                                                                                                                 12:
       5: \mathbf{z}_i := \bot
                                                                                                                 13: if \tilde{\mathbf{u}}_1 = \mathbf{u}_1 then return 1
       6: while \mathbf{z}_i := \bot \mathbf{do}
                     \mathbf{y}_i \leftarrow \mathcal{D}^{\ell+k}
                      \mathbf{u}_i := \bar{\mathbf{A}} \mathbf{y}_i \mod q
       9:
                      \tilde{\mathbf{u}}_i := \tilde{\mathbf{u}}_{i-1} + \mathbf{u}_i \bmod q
                      \overline{c_i} := \mathsf{H}(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})
      10:
                      \mathbf{z}_i := \overline{c_i \cdot \mathbf{s}_i} + \mathbf{y}_i
      11:
                      \mathbf{z}_i := \mathsf{RejSamp}(\mathbf{z}_i, c_i \cdot \mathbf{s}_i)
      12:
     13: \sigma_i := (\tilde{\mathbf{u}}_i, \mathbf{z}_1, \dots, \mathbf{z}_i)
     14: return \sigma_i
```

Remark 3.1. Whereas it seems to be hard to give unconditionally provable lower bounds for p_{inv} (at least in our parameter setting), it is possible to bound it assuming the hardness of M-LWE (which is also used in the security proof of FSwA-S). Let p_{R_q} denote the probability that an element of R_q sampled uniformly at random is invertible. There exist exact formulas to express this number, depending on the splitting behavior of the ideal generated by q in the ring R. For the fully splitting case, i.e., $q=1 \mod 2n$, it yields $p_{R_q}=(1-1/q)^n$, see for instance [LPR13, Claim 2.25]. Assuming the hardness of M-LWE, it yields $p_{\text{inv}}=p_{R_q}+\text{negl}(\lambda)$. If not, an adversary against M-LWE could simply test a given instance for invertibility.

Remark 3.2. In its current presentation, all secret-public key pairs are using the same matrix A. However, this is not required in our construction, and actually every party could use their own matrix. In that case, the public key would need to contain not only t, but also A, requiring larger storage. This, in turn, could be reduced by computing the matrix via some small seed and a pseudorandom function (as done in Dilithium). We highlight that using different matrices for different parties is not possible in all other proposed (non-interactive or interactive) aggregate signature schemes following the Fiat-Shamir with aborts paradigm. As they compute a linear combination of (parts of) the single signatures, every party must use the same A. Using the same matrix A for every party leads to an instance of multi-secret M-LWE and its security is implied by standard M-LWE via some simple hybrid argument [Mic18, Lemma 8].

Lemma 3.3 (Correctness). The scheme FSwA-SAS = (Setup, Gen, SeqSign, SeqVerify) as specified in Algorithm 2 is correct.

Proof. We inductively show that, if an *i*-th so-far signature $\sigma_i = (\tilde{\mathbf{u}}_i, \mathbf{z}_1, \dots, \mathbf{z}_i)$ with $1 \leq i < N$ is correct, the (i+1)-th signature $\sigma_{i+1} = (\tilde{\mathbf{u}}_{i+1}, \mathbf{z}_1, \dots, \mathbf{z}_{i+1})$ is also correct. As \mathbf{z}_{i+1} has been correctly computed by the (i+1)-th signer, it yields $\tilde{\mathbf{A}} \cdot \mathbf{z}_{i+1} - \mathbf{t}_{i+1} \cdot c_{i+1} = \mathbf{u}_{i+1}$. Hence, $\tilde{\mathbf{u}}_i$ can be recovered via $\tilde{\mathbf{u}}_{i+1} - \mathbf{u}_{i+1}$ and thus $\sigma_i = (\tilde{\mathbf{u}}_i, \mathbf{z}_1, \dots, \mathbf{z}_i)$ verifies by the induction hypothesis. Now,

let's consider the base case i=1. It yields $\|\mathbf{z}_1\|_{\infty} \leq B$ and $\mathbf{\bar{A}} \cdot \mathbf{z}_1 = \mathbf{t}_1 \cdot c_1 + \tilde{\mathbf{u}}_1$ because of the linearity of matrix-vector multiplication over R_q .

3.2 Security Proof

We now prove the FH-UF-CMA security (as in Definition 2.9) of Algorithm 2. For completeness, we also sketch in Appendix A the security of our scheme in a new model that has been introduced in [CZ22], which we call the *partial-signature history-free* security model.

Theorem 3.4 (FH-UF-CMA security). Let $k, \ell, n, q, \eta, \gamma, l \in \mathbb{N}$ such that n is a power of 2, q is prime and Eq. 1 is fulfilled. Let p_{inv} be the probability that $\bar{\mathbf{A}}\mathbf{s}$ has at least one invertible coefficient over R_q , where \mathbf{s} is uniformly sampled from $S_\eta^{\ell+k}$ and $\bar{\mathbf{A}} = [\mathbf{A}|\mathbf{I}_k]$ with \mathbf{A} is uniformly sampled from $R_q^{\mathsf{k}\times\ell}$, respectively. If the signature scheme FSwA-S with message space $M = \{0,1\}^l$, as described in Algorithm 1, is UF-CMA secure, then is the sequential aggregate signature FSwA-SAS, as described in Algorithm 2, FH-UF-CMA secure. Concretely, for any adversary A against FH-UF-CMA security that makes at most Q_h queries to the random oracle H, Q_s queries to the OSeqSign oracle and outputs a forgery with a history of length N, there exists an adversary B against UF-CMA security such that

$$\frac{\mathbf{Adv}_{\mathsf{FSwA-SAS}}^{\mathsf{FH-UF-CMA}}(\mathcal{A}) \leq \frac{\mathbf{Adv}_{\mathsf{FSwA-S}}^{\mathsf{UF-CMA}}(\mathcal{B})}{p_{\mathsf{inv}}} + O\left(\frac{Q_s(Q_h + Q_s)}{q^{nk/2}}\right) + \frac{(Q_h + Q_s + 1)^2}{|\mathsf{Ch}|} + \frac{Q_s(2Q_h + 1)}{2^l},$$

and $\text{Time}(\mathcal{B}) = \text{Time}(\mathcal{A}) + O((N + Q_h)k\ell t_{\text{pmul}})$, where t_{pmul} is the time of polynomial multiplication in R_q .

Proof. We first sketch the high level ideas of the reduction \mathcal{B} . The complete description of \mathcal{B} is found in Alg. 3. The random oracle and the signing oracle in the FH-UF-CMA game (resp. UF-CMA game) are denoted by H and OSeqSign (resp. H' and OSign). On receiving the public parameter A and the challenge public key \mathbf{t}^* , \mathcal{B} checks that $\mathbf{t}^* \in R_q^k$ contains at least one invertible element. If so, \mathcal{B} forwards $(\mathbf{A}, \mathbf{t}^*)$ to \mathcal{A} .

OSeqSign replies to queries by asking OSign for a signature on uniformly chosen m and programs H such that it outputs c returned by the outer random oracle H'. Here we cannot just forward m_i to OSign, because it might be that a forgery submitted by \mathcal{A} later reuses the same m_i . Then submitting a forgery w.r.t. m_i is not valid in the UF-CMA game, causing \mathcal{B} to lose.

At the core of reduction is simulation of responses to H queries. Suppose the key list L_N as part of the forgery tuple contains (\mathbf{t}_i, m_i) such that $\mathbf{t}_i = \mathbf{t}^*$. Then \mathcal{B} must have extracted the corresponding \mathbf{u}_i and forwarded \mathbf{u}_i to H' together with a random message m, so that $(m, (\mathbf{u}_i, \mathbf{z}_i))$ qualifies as a valid forgery in the UF-CMA game. This extraction operation crucially makes use of \mathbf{z}_{i-1} when $(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})$ is queried to H. Intuitively, \mathbf{z}_{i-1} serves as a look-up key to obtain the previous aggregated $\tilde{\mathbf{u}}_{i-1}$, which allows \mathcal{B} to extract $\mathbf{u}_i = \tilde{\mathbf{u}}_i - \tilde{\mathbf{u}}_{i-1}$.

In more detail, starting from the original FH-UF-CMA game, we construct several hybrid games towards the one used by the final reduction \mathcal{B} . We denote by $\Pr[G_i(\mathcal{A})]$ the probability that $G_i(\mathcal{A})$ halts with output 1.

- G_0 This game is identical to the FH-UF-CMA game. At the beginning, the game initializes an empty key-value look-up table HT. Upon receiving a query to the random oracle H with input X, it returns $\operatorname{HT}[X]$ if the table entry is non-empty; otherwise, it samples uniform $c \in \operatorname{Ch}$, sets $\operatorname{HT}[X] := c$, and returns c. It holds that $\operatorname{Pr}[G_0(\mathcal{A})] = \operatorname{Adv}_{\mathsf{FSwA-SAS}}^{\mathsf{FH-UF-CMA}}(\mathcal{A})$.
- G₁ This game is identical to G₀, except that OSeqSign samples uniform $c_i \in \mathsf{Ch}$ instead of calling $c_i = \mathsf{H}(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})$ after $\tilde{\mathbf{u}}_i$ is computed, and that it programs the RO table $\mathsf{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] := c_i$ as soon as the rejection sampling step succeeds; if $\mathsf{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ is already set, the game aborts by setting $\mathsf{bad}_{\mathsf{ucol}} = \mathsf{true}$. It holds that $|\mathsf{Pr}[\mathsf{G}_0(\mathcal{A})] \mathsf{Pr}[\mathsf{G}_1(\mathcal{A})]| \leq \mathsf{Pr}[\mathsf{bad}_{\mathsf{ucol}}]$.
- G₂ This game is identical to G₁, except that responses to random oracle queries $\mathsf{H}(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})$ are simulated as follows. Initialize an empty key-value look-up table ZT. If i=1 or there exists some $X:=(\tilde{\mathbf{u}}_{i-1}, L_{i-1}, \mathbf{z}_{i-2})$ such that $\mathrm{ZT}[X]=\bar{\mathbf{A}}\mathbf{z}_{i-1} \bmod q$, then extract $\mathbf{u}_i:=\tilde{\mathbf{u}}_i-\tilde{\mathbf{u}}_{i-1}$, sample uniform $c_i\in\mathsf{Ch}$, and set $\mathrm{ZT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]:=\mathbf{u}_i+c_i\mathbf{t}_i$. If there already exists some entry $X':=(\tilde{\mathbf{u}}_{i-1}, L_{i-1}, \mathbf{z}_{i-2})$ such that $\mathrm{ZT}[X']=\mathbf{u}_i+c_i\mathbf{t}_i$, the game aborts by setting $\mathsf{bad}_{\mathsf{zcol}}=\mathsf{true}$. It holds that $|\Pr[\mathsf{G}_1(\mathcal{A})]-\Pr[\mathsf{G}_2(\mathcal{A})]|\leq \Pr[\mathsf{bad}_{\mathsf{zcol}}]$.

- G₃ This game is identical to G₂, except that OSeqSign and H proceed as follows. The game initializes an empty set \mathcal{M} and key-value look-up table MT. Whenever OSeqSign receives a query, it internally samples a uniform message $m \in \mathcal{M}$ and adds m to \mathcal{M} . Whenever H receives a query with input $(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})$ and manages to extract \mathbf{u}_i as above, it samples a uniform message $m \in \mathcal{M}$ and aborts by setting $\mathsf{bad}_{\mathsf{mcol}} = \mathsf{true}$ if $m \in \mathcal{M}$. Else, it sets $\mathsf{MT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] = m$. It holds that $|\Pr[\mathsf{G}_2(\mathcal{A})] \Pr[\mathsf{G}_3(\mathcal{A})]| \leq \Pr[\mathsf{bad}_{\mathsf{mcol}}]$.
- G_4 This game is identical to G_3 , except that it performs the following checks against the ZT entries after the adversary outputs a valid signature-history pair $(L_N, (\tilde{\mathbf{u}}_N, \mathbf{z}_1, \dots, \mathbf{z}_N))$ as follows. Let $\tilde{\mathbf{u}}_{N-1}, \dots, \tilde{\mathbf{u}}_1$ be as derived during the execution of SeqVerify. If for some $i \in [N]$ the entry $\mathrm{ZT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ is undefined, the game halts by setting $\mathsf{bad}_\mathsf{ord} = \mathsf{true}$. It holds that $|\Pr[G_3(\mathcal{A})] \Pr[G_4(\mathcal{A})]| \leq \Pr[\mathsf{bad}_\mathsf{ord}]$.
- \mathcal{B} Given an adversary \mathcal{A} winning G_4 , the reduction \mathcal{B} described in Alg. 3 is obtained as follows. Upon receiving a query to OSeqSign, \mathcal{B} makes a query to OSign of the UF-CMA game with a uniform message $m \in \mathcal{M}$, receives \mathbf{u}_i and \mathbf{z}_i , and programs HT using challenge c_i output by the outer random oracle $H'(\mathbf{u}_i, \mathbf{t}^*, m)$. Moreover, H obtains fresh challenge c_i for $\mathbf{t}_i = \mathbf{t}^*$ by querying the outer random oracle $H'(\mathbf{u}_i, \mathbf{t}^*, m)$ if it succeeds in extracting $\mathbf{u}_i = \tilde{\mathbf{u}}_i \tilde{\mathbf{u}}_{i-1}$. Since \mathcal{A} is guaranteed to receive an invertible challenge public key in \mathcal{B} , the view of \mathcal{A} is identical to that of G_4 .

We now show that, as long as none of the bad events happen, \mathcal{B} is guaranteed to output a message-signature pair $(m, (\mathbf{u}_{i^*}, \mathbf{z}_{i^*}))$ that gets accepted in the UF-CMA game, i.e., $\|\mathbf{z}_{i^*}\|_{\infty} \leq B$ and $\mathbf{u}_{i^*} = \bar{\mathbf{A}}\mathbf{z}_{i^*} - c\mathbf{t}^*$ mod q where $c = \mathsf{H}'(\mathbf{u}_{i^*}, \mathbf{t}^*, m)$. The former condition is immediate from the verification condition of SeqVerify. To argue the latter, notice that we have $c = \mathsf{H}'(\mathbf{u}_{i^*}, \mathbf{t}^*, m) = \mathsf{HT}[\tilde{\mathbf{u}}_{i^*}, L_{i^*}, \mathbf{z}_{i^*-1}] = c_{i^*}$ as long as the RO entries $\mathsf{HT}[\tilde{\mathbf{u}}_1, L_1, \mathbf{z}_0], \ldots, \mathsf{HT}[\tilde{\mathbf{u}}_N, L_N, \mathbf{z}_{N-1}]$ have been set in the right order and thus $\mathbf{u}_{i^*} = \tilde{\mathbf{u}}_{i^*} - \tilde{\mathbf{u}}_{i^*-1}$ is extracted during the invocation of $\mathsf{H}(\tilde{\mathbf{u}}_{i^*}, L_{i^*}, \mathbf{z}_{i^*-1})$. The following lemma indeed assures that such queries have been made in the right order as long as $\mathsf{bad}_{\mathsf{zcol}} = \mathsf{bad}_{\mathsf{ord}} = \mathsf{false}$.

Lemma 3.5. Let $\sigma_N = (\tilde{\mathbf{u}}_N, \mathbf{z}_1, \dots, \mathbf{z}_N)$ and $L_N = (\mathbf{t}_1, m_1)||\dots||(\mathbf{t}_N, m_N)$ a valid signature-history pair that \mathcal{B} received from \mathcal{A} . Let $\tilde{\mathbf{u}}_1, \dots, \tilde{\mathbf{u}}_{N-1}$ be as derived in SeqVerify run by \mathcal{B} . Suppose $\mathsf{bad}_{\mathsf{zcol}} = \mathsf{false}$. Then for $i \in [N-1]$, the random oracle entry $\mathsf{HT}[\tilde{\mathbf{u}}_{i+1}, L_{i+1}, \mathbf{z}_i]$ had been set after $\mathsf{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ was set if and only if $\mathsf{bad}_{\mathsf{ord}} = \mathsf{false}$.

Proof. "Only if" We first argue that if the oracle entries $\operatorname{HT}[\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0],\ldots,\operatorname{HT}[\tilde{\mathbf{u}}_N,L_N,\mathbf{z}_{N-1}]$ have been set in this order and $\operatorname{\mathsf{bad}}_{\operatorname{\mathsf{zcol}}} = \operatorname{\mathsf{false}}$, the corresponding ZT entries are all non-empty and thus $\operatorname{\mathsf{bad}}_{\operatorname{\mathsf{ord}}}$ must be false. Suppose this statement holds for $1 \leq i \leq j$, i.e., $\operatorname{ZT}[\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0],\ldots,\operatorname{ZT}[\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1}]$ are non-empty. Due to the verification condition it must be that $\tilde{\mathbf{u}}_{j-1} = \tilde{\mathbf{u}}_j - \bar{\mathbf{A}}\mathbf{z}_j + c_j\mathbf{t}_j$ and thus $\bar{\mathbf{A}}\mathbf{z}_j = \tilde{\mathbf{u}}_j - \tilde{\mathbf{u}}_{j-1} + c_j + \mathbf{t}_j$. Because we assumed that $\operatorname{HT}[\tilde{\mathbf{u}}_{j-1},L_{j-1},\mathbf{z}_{j-2}]$ and $\operatorname{HT}[\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1}]$ are set in this order, the invocation of $\operatorname{\mathsf{H}}(\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1})$ must have extracted $\mathbf{u}_j = \tilde{\mathbf{u}}_j - \tilde{\mathbf{u}}_{j-1}$ and have set $\operatorname{ZT}[\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1}] = \mathbf{u}_j + c_j\mathbf{t}_j$. Note that, since $\operatorname{\mathsf{bad}}_{\operatorname{\mathsf{zcol}}} = \operatorname{\mathsf{false}}$, there is no other entry in ZT that records the same value as $\mathbf{u}_j + c_j\mathbf{t}_j$. Thus, when $(\tilde{\mathbf{u}}_{j+1},L_{j+1},\mathbf{z}_j)$ is queried, $\operatorname{\mathsf{H}}$ can uniquely find a tuple $(\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1})$ such that $\operatorname{ZT}[\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1}] = \bar{\mathbf{A}}\mathbf{z}_j$ and then set $\operatorname{ZT}[\tilde{\mathbf{u}}_{j+1},L_{j+1},\mathbf{z}_j] = \tilde{\mathbf{u}}_{j+1} - \tilde{\mathbf{u}}_j + c_{j+1}\mathbf{t}_{j+1}$. It is easy to see that the base case j=1 is true: when $(\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0)$ is queried, the invocation of $\operatorname{\mathsf{H}}$ always sets $\operatorname{ZT}[\tilde{\mathbf{u}}_j,L_1,\mathbf{z}_0] = \mathbf{u}_1 + c_1\mathbf{t}_1$ where $\mathbf{u}_1 = \tilde{\mathbf{u}}_1$.

"If" We give a proof by induction. As an induction hypothesis, we assume that for $i=1,\ldots,j-1$ the random oracle entry $\operatorname{HT}[\tilde{\mathbf{u}}_{i+1},L_{i+1},\mathbf{z}_i]$ had been set after $\operatorname{HT}[\tilde{\mathbf{u}}_i,L_i,\mathbf{z}_{i-1}]$ was set whenever $\operatorname{\mathsf{bad}}_{\mathsf{zcol}} = \operatorname{\mathsf{bad}}_{\mathsf{ord}} = \mathsf{false}$. Now suppose, for a contradiction, that $\operatorname{HT}[\tilde{\mathbf{u}}_{i+1},L_{j+1},\mathbf{z}_j]$ was set before $\operatorname{HT}[\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1}]$ while $\operatorname{\mathsf{bad}}_{\mathsf{zcol}} = \operatorname{\mathsf{bad}}_{\mathsf{ord}} = \mathsf{false}$. Because $\operatorname{\mathsf{bad}}_{\mathsf{ord}} = \mathsf{false}$, the entry $\operatorname{ZT}[\tilde{\mathbf{u}}_i,L_i,\mathbf{z}_{i-1}]$ is non-empty for all $i\in[N]$. When H is queried with input $(\tilde{\mathbf{u}}_{j+1},L_{j+1},\mathbf{z}_j)$, since the corresponding entry in ZT is non-empty, it must be that there exists some $X':=(\tilde{\mathbf{u}}'_j,L_j,\mathbf{z}'_{j-1})\neq (\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1})$ such that $c'_j:=\operatorname{HT}[X']$ and $\operatorname{ZT}[X']=\bar{\mathbf{A}}\mathbf{z}_j$ are already set. This implies that X' has been queried to H before and that $\bar{\mathbf{A}}\mathbf{z}_j=\mathbf{u}'_j+c'_j\mathbf{t}_j \bmod q$, where \mathbf{u}'_j is some value extracted inside $\operatorname{H}(X')$. On the other hand, due to the verification condition it also holds that $\bar{\mathbf{A}}\mathbf{z}_j=\mathbf{u}_j+c_j\mathbf{t}_j \bmod q$, where $c_j=\operatorname{HT}[\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1}]$ and $\mathbf{u}_j=\tilde{\mathbf{u}}_j-\tilde{\mathbf{u}}_{j-1}$. Here, \mathbf{u}_j is the value extracted when $(\mathbf{u}_j,L_j,\mathbf{z}_{j-1})$ is queried to H for the first time, because due to the induction hypothesis $\operatorname{HT}[\tilde{\mathbf{u}}_{j-1},L_{j-1},\mathbf{z}_{j-2}]$ had been already set at this point. However, this implies that $\operatorname{\mathsf{bad}}_{\mathsf{zcol}}$ is set when $(\tilde{\mathbf{u}}_j,L_j,\mathbf{z}_{j-1})$ is queried to H, contradicting the assumption that $\operatorname{\mathsf{bad}}_{\mathsf{zcol}} = \operatorname{\mathsf{false}}$.

Let us prove the base case j=2 in a similar manner. Suppose $\operatorname{HT}[\tilde{\mathbf{u}}_2,L_2,\mathbf{z}_1]$ was set before $\operatorname{HT}[\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0]$ while $\operatorname{\mathsf{bad}}_{\mathsf{zcol}}=\operatorname{\mathsf{bad}}_{\mathsf{ord}}=\mathsf{false}$. When H is queried with input $(\tilde{\mathbf{u}}_2,L_2,\mathbf{z}_1)$, since the corresponding entry in ZT is non-empty, it must be that there exists some $X':=(\tilde{\mathbf{u}}_1',L_1,\mathbf{z}_0)\neq (\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0)$ such that $c_1':=\operatorname{HT}[X']$ and $\operatorname{ZT}[X']=\bar{\mathbf{A}}\mathbf{z}_1$ are already set. This implies that X' has been queried to H before and that $\bar{\mathbf{A}}\mathbf{z}_1=\tilde{\mathbf{u}}_1'+c_1'\mathbf{t}_1 \bmod q$. On the other hand, due to the verification condition it also holds that $\bar{\mathbf{A}}\mathbf{z}_1=\tilde{\mathbf{u}}_1+c_1\mathbf{t}_1 \bmod q$, where $c_1=\operatorname{HT}[\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0]$. However, this implies that $\operatorname{\mathsf{bad}}_{\mathsf{zcol}}$ is set when $(\tilde{\mathbf{u}}_1,L_1,\mathbf{z}_0)$ is queried to H, contradicting the assumption that $\operatorname{\mathsf{bad}}_{\mathsf{zcol}}=\mathsf{false}$.

All in all, unless \mathcal{B} sets $\mathsf{bad}_{\mathsf{inv}} = \mathsf{true}$, \mathcal{B} wins the UF-CMA game if and only if G_4 outputs 1. In other words,

$$\begin{split} \mathbf{Adv}_{\mathsf{FSwA-S}}^{\mathsf{UF-CMA}}(\mathcal{B}) &= (1 - \Pr[\mathsf{bad}_{\mathsf{inv}}]) \cdot \Pr[G_4] \\ &\geq p_{\mathsf{inv}} \cdot \left(\mathbf{Adv}_{\mathsf{FSwA-SAS}}^{\mathsf{FH-UF-CMA}}(\mathcal{A}) - \Pr[\mathsf{bad}_{\mathsf{ucol}}] - \Pr[\mathsf{bad}_{\mathsf{zcol}}] - \Pr[\mathsf{bad}_{\mathsf{mcol}}] - \Pr[\mathsf{bad}_{\mathsf{ord}}] \right). \end{split}$$

The running time of \mathcal{B} is at most the running time of \mathcal{A} plus the time it takes for running verification operations and handling random oracle queries. The former takes $O(Nk\ell t_{\mathsf{pmul}})$ because each iteration of the for-loop involves matrix-vector multiplication $\bar{\mathbf{A}}\mathbf{z}_i$ (ignoring the run-time for polynomial addition as it's much smaller than multiplication). The latter takes $O(Q_hk\ell t_{\mathsf{pmul}})$ because \mathcal{B} carries out matrix-vector multiplication $\bar{\mathbf{A}}\mathbf{z}_{i-1}$ for each query to the RO H.

In the following, we provide a concrete bound for each bad event.

Bounding $\Pr[\mathsf{bad}_{\mathsf{mcol}}]$ Assuming that the adversary makes at most Q_s queries to $\mathsf{OSeqSign}$, there are at most Q_s distinct values in \mathcal{M} and Q_h distinct values in MT , respectively. The $\mathsf{bad}_{\mathsf{mcol}}$ flag is potentially set due to two different causes: (1) H internally samples m that is already recorded in \mathcal{M} and thus the corresponding entry is not stored in MT , or (2) $\mathsf{OSeqSign}$ internally samples m that is already recorded in MT and thus the corresponding entry gets removed. The probability that a randomly sampled m inside H collides with one of the values in \mathcal{M} is at most $Q_s/|\mathcal{M}| = Q_s/2^l$. Since at most Q_h queries to H are made by \mathcal{A} , the probability that case (1) occurs during such queries is at most $Q_s \cdot Q_h/2^l$. If for some $i \in I$ a tuple $(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})$ is queried to H during the invocation of $\mathsf{SeqVerify}$ for the first time, in order to cause $\mathsf{bad}_{\mathsf{mcol}} = \mathsf{true}$, for all such i independently sampled m must be in \mathcal{M} . Thus, the probability that case (1) occurs during such queries is at most $Q_s/2^l$. The probability that a randomly sample m inside $\mathsf{OSeqSign}$ collides with one of the values in MT is at most $Q_h/2^l$. Since at most Q_s queries to $\mathsf{OSeqSign}$ are made, the probability that case (2) occurs is at most $Q_s \cdot Q_h/2^l$. Overall, we get $\mathsf{Pr}[\mathsf{bad}_{\mathsf{mcol}}] \leq Q_s \cdot (2Q_h + 1)/2^l$.

Bounding $\Pr[\mathsf{bad}_{\mathsf{ucol}}]$ Since there are at most $Q_h + Q_s$ values in HT and \mathbf{u} is generated by the signing algorithm Sign from FSwA-S, for each query to OSeqSign the probability that the flag $\mathsf{bad}_{\mathsf{ucol}}$ is set is at most

$$\max_{\mathbf{u}} \Pr \left[\mathbf{u} = \bar{\mathbf{A}} \mathbf{y} \bmod q : \mathbf{y} \stackrel{\$}{\leftarrow} \mathscr{D}^{\ell+k} \right]. \tag{2}$$

Equation 2 can be upper bounded by Lemma 2.2, using the probability preservation property of the Rényi divergence. It yields

$$\max_{\mathbf{u}} \Pr\left[\mathbf{u} = \bar{\mathbf{A}}\mathbf{y} \bmod q : \mathbf{y} \stackrel{\$}{\leftarrow} \mathscr{D}^{\ell+k}\right]$$

$$\leq \sqrt{\max_{\mathbf{u}} \Pr\left[\mathbf{u} = \mathbf{v} \bmod q : \mathbf{v} \stackrel{\$}{\leftarrow} R_q^k\right]} \cdot \operatorname{RD}_2((\bar{\mathbf{A}}, \bar{\mathbf{A}}\mathbf{y}) \| (\bar{\mathbf{A}}, \mathbf{v})).$$

The probability in the second inequality is given by $\sqrt{1/|R_q^k|} = q^{-nk/2}$. By Lem. 2.2, using the data processing inequality, it holds $\mathrm{RD}_2((\bar{\mathbf{A}},\bar{\mathbf{A}}\mathbf{y})\|(\bar{\mathbf{A}},\mathbf{v})) \leq \mathrm{RD}_2((\mathbf{A},\mathbf{A}\mathbf{y}')\|(\mathbf{A},\mathbf{v}))$, where $\mathbf{y}' \stackrel{\$}{\leftarrow} \mathscr{D}^\ell$. By Lem. 2.3, the latter is bounded above by a constant if Eq. 1 is fulfilled. Since OSeqSign receives at most Q_s queries, overall, we obtain $\Pr[\mathsf{bad}_{\mathsf{ucol}}] \leq O(Q_s(Q_h + Q_s)/q^{nk/2})$.

Bounding $\Pr[\mathsf{bad}_{\mathsf{zcol}}]$ Fix an existing entry in ZT of the form $\tilde{\mathbf{u}}_i' + c_i'\mathbf{t}_i \mod q$. Then the probability that $\mathbf{u}_i + c_i\mathbf{t}_i \mod q$ hits such an entry is

$$\Pr_{\substack{\mathbf{c}_i \overset{\$}\leftarrow \mathsf{Ch}}} \left[\mathbf{u}_i + c_i \mathbf{t}_i = \tilde{\mathbf{u}}_i' + c_i' \mathbf{t}_i \bmod q \right] = \Pr_{\substack{\mathbf{c}_i \overset{\$}\leftarrow \mathsf{Ch}}} \left[c_i \mathbf{t}_i = \tilde{\mathbf{u}}_i' + c_i' \mathbf{t}_i - \mathbf{u}_i \bmod q \right].$$

Algorithm 3: Reduction to UF-CMA security of FSwA-S

The random oracle in the UF-CMA game is denoted by H'. The sign oracle in the UF-CMA game is denoted by OSign. Let $\tilde{\mathbf{u}}_0 = 0$. Without loss of generality, \mathcal{A} queries H with input public keys $\mathbf{t}_1, \dots, \mathbf{t}_i$ all of which contain at least one invertible element, because otherwise such keys will be rejected by the verification algorithm anyway. All flags are initially set to false.

```
\mathcal{B}^{\mathsf{OSign},\mathsf{H}'}(\bar{\mathbf{A}},\mathbf{t}^*)
                                                                                                                              \mathsf{H}(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})
       1: Q := \emptyset; \mathcal{M} := \emptyset
                                                                                                                                     1: if HT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] \neq \bot then
      2: if \mathbf{t}^* has no invertible element then
                                                                                                                                                  return HT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]
                     \mathsf{bad}_{\mathsf{inv}} := \mathsf{true}
                                                                                                                                           (\mathbf{t}_1, m_1)||\ldots||(\mathbf{t}_i, m_i) := L_i
           (\sigma_N, L_N) \leftarrow \mathcal{A}^{\mathsf{OSeqSign}, \mathsf{H}}(\bar{\mathbf{A}}, \mathbf{t}^*)
                                                                                                                                     4: if i = 1 or \exists X := (\tilde{\mathbf{u}}_{i-1}, L_{i-1}, \mathbf{z}_{i-2}) such
           (\mathbf{t}_1, m_1)||\ldots||(\mathbf{t}_N, m_N) := L_N
                                                                                                                                           that HT[X] \neq \bot \land ZT[X] = \bar{\mathbf{A}}\mathbf{z}_{i-1} \mod q then
           (\tilde{\mathbf{u}}_N, \mathbf{z}_1 \dots, \mathbf{z}_N) := \sigma_N
                                                                                                                                     5:
                                                                                                                                                    \mathbf{u}_i := \tilde{\mathbf{u}}_i - \tilde{\mathbf{u}}_{i-1} \bmod q
      7: I := \{ i \in [N] : \mathbf{t}_i = \mathbf{t}^* \land (m_i, L_{i-1}) \notin \mathcal{Q} \}
                                                                                                                                     6:
                                                                                                                                                  if \mathbf{t}_i = \mathbf{t}^* then
      8: if SeqVerify(\sigma_N, L_N) = 1 \land |I| \neq 0 then
                                                                                                                                                           m \stackrel{\mathfrak{s}}{\leftarrow} M
                                                                                                                                     7:
                   Derive \tilde{\mathbf{u}}_{N-1}, \dots, \tilde{\mathbf{u}}_1 as in SeqVerify
                                                                                                                                                           c_i := \mathsf{H}'(\mathbf{u}_i, m, \mathbf{t}^*)
                                                                                                                                     8:
                   if \exists i \in [N] such that ZT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] = \bot then
    10:
                                                                                                                                                         if m \notin \mathcal{M} then
                                                                                                                                     9:
    11:
                            bad_{ord} := true
                                                                                                                                                                  MT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] := m
                                                                                                                                   10:
                   if \exists i^* \in I such that m := MT[\tilde{\mathbf{u}}_{i^*}, L_{i^*}, \mathbf{z}_{i^*-1}] \neq
                                                                                                                                                  else
                                                                                                                                   11:
            \perp then
                                                                                                                                                           c_i \stackrel{\$}{\leftarrow} \mathsf{Ch}
                                                                                                                                   12:
    13:
                            \mathbf{u}_{i^*} := \tilde{\mathbf{u}}_{i^*} - \tilde{\mathbf{u}}_{i^*-1}
                                                                                                                                   13:
                                                                                                                                                  if \exists X' := (\tilde{\mathbf{u}}_i', L_i, \mathbf{z}_{i-1}') such that \mathrm{ZT}[X'] = \mathbf{u}_i +
                           return (m, \mathbf{u}_{i^*}, \mathbf{z}_{i^*})
    14:
                                                                                                                                           c_i \mathbf{t}_i \bmod q then
    15:
                   else
                                                                                                                                   14:
                                                                                                                                                           bad_{zcol} := true
                            bad_{mcol} := true
                                                                                                                                                    ZT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] := \mathbf{u}_i + c_i \mathbf{t}_i \mod q
                                                                                                                                   15:
\mathsf{OSeqSign}(m_i, L_{i-1}, \sigma_{i-1})
                                                                                                                                   16: else
                                                                                                                                                    c_i \stackrel{\$}{\leftarrow} \mathsf{Ch}
      1: Q := Q \cup \{(m_i, L_{i-1})\}
                                                                                                                                   17:
      2: (\tilde{\mathbf{u}}_{i-1}, \mathbf{z}_1, \dots, \mathbf{z}_{i-1}) := \sigma_{i-1}
                                                                                                                                           \mathrm{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] := c_i
      3: m \stackrel{\$}{\leftarrow} M
                                                                                                                                   19: return c_i
      4: if \exists X such that MT[X] = m then
                     \mathrm{MT}[X] := \bot
      6: \mathcal{M} := \mathcal{M} \cup \{m\}
             (\mathbf{u}, \mathbf{z}) \leftarrow \mathsf{OSign}(m)
      8: c := \mathsf{H}'(\mathbf{u}, m, \mathbf{t}^*)
     9: L_i := L_{i-1}||(\mathbf{t}^*, m_i)||
    10: \tilde{\mathbf{u}}_i := \tilde{\mathbf{u}}_{i-1} + \mathbf{u} \mod q
    11: \mathbf{z}_i := \mathbf{z}
    12: if HT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] \neq \bot then
    13:
                     bad_{ucol} := true
    14: else
                     HT[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}] := c
    16: \sigma_i := (\tilde{\mathbf{u}}_i, \mathbf{z}_1, \dots, \mathbf{z}_i)
    17: return \sigma_i
```

Since at least one coefficient of \mathbf{t}_i is invertible, the above probability is bounded by $1/|\mathsf{Ch}|$. Let Q_i be the number of entries in HT indexed by a tuple containing a history L of size i. Then we have $Q_h + Q_s = \sum_{i=1}^N Q_i$. Because H receives at most $Q_i + 1$ queries for each (where "+1" comes from the fact that an additional query is made from inside SeqVerify), by the union bound, we have that $\Pr[\mathsf{bad}_{\mathsf{zcol}}] \leq \sum_{i=1}^N Q_i(Q_i + 1)/(2|\mathsf{Ch}|)$.

Bounding $\Pr[\mathsf{bad}_{\mathsf{ord}}]$ Due to Lemma 3.5, " $\mathsf{bad}_{\mathsf{ord}} = \mathsf{true}$ while $\mathsf{bad}_{\mathsf{zcol}} = \mathsf{false}$ " implies that there exists some $i \in [N-1]$, such that the entry $\operatorname{HT}[\tilde{\mathbf{u}}_{i+1}, L_{i+1}, \mathbf{z}_i]$ was set before $\operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$, where $(\tilde{\mathbf{u}}_N, (\mathbf{z}_1, \ldots, \mathbf{z}_N))$ and L_N are signature-history pair output by \mathcal{A} at the end of the game. We argue that this event occurs with negligible probability if the verification condition is satisfied. The event potentially occurs in two ways: (1) for some $i \in [N-1]$, $\operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ is set for the first time during the invocation of SeqVerify, and (2) for some $i \in [N-1]$, $\operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ was set for the first time before the invocation of SeqVerify, but after $\operatorname{HT}[\tilde{\mathbf{u}}_{i+1}, L_{i+1}, \mathbf{z}_i]$ was set.

To bound case (1), it is sufficient to prove the following statement inductively: if none of $\operatorname{HT}[\tilde{\mathbf{u}}_1, L_1, \mathbf{z}_0], \ldots, \operatorname{HT}[\tilde{\mathbf{u}}_{i-1}, L_{i-1}, \mathbf{z}_{i-2}]$ have been set for the first time during the invocation of SeqVerify, the probability that $\operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ is set for the first time while running SeqVerify and that the verification condition is satisfied, is at most $Q_{i-1}/|\operatorname{Ch}|$, where Q_i 's are defined as above and let $Q_0 = 1$ for convenience. Since $\tilde{\mathbf{u}}_{i-1}, \tilde{\mathbf{u}}_i, \mathbf{z}_i, \mathbf{t}_i$ have been already fixed at the moment when $(\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1})$ is queried to H inside SeqVerify, the probability that the signature gets accepted is

$$\Pr_{\substack{c_i \overset{\leftarrow}{\leftarrow} \mathsf{Ch}}} \left[\tilde{\mathbf{u}}_{i-1} = \tilde{\mathbf{u}}_i - (\bar{\mathbf{A}} \mathbf{z}_i - c_i \mathbf{t}_i) \bmod q \right],$$

which is at most $1/|\mathsf{Ch}|$. Because there are at most Q_{i-1} entries HT indexed by a tuple containing L_i and thus at most Q_{i-1} different values for $\tilde{\mathbf{u}}_{i-1}$ exist, by the union bound, the probability that case (1) happens is at most $Q_{i-1}/|\mathsf{Ch}|$. The base case is true: if i=1, since $\tilde{\mathbf{u}}_1, \mathbf{z}_1, \mathbf{t}_1$ have been already fixed at the moment when $(\tilde{\mathbf{u}}_1, L_1, \mathbf{z}_0)$ is queried to H inside SeqVerify, the probability that the signature gets accepted is

$$\Pr_{\substack{c_1 \stackrel{\$}{\leftarrow} \mathsf{Ch}}} \left[\tilde{\mathbf{u}}_1 = \bar{\mathbf{A}} \mathbf{z}_1 - c_1 \mathbf{t}_1 \bmod q \right],$$

which is at most 1/|Ch|.

To bound case (2), it is sufficient to prove the following statement inductively: if $\operatorname{HT}[\tilde{\mathbf{u}}_1, L_1, \mathbf{z}_0], \ldots, \operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ have been set in this order before the invocation of SeqVerify, the probability that $\operatorname{HT}[\tilde{\mathbf{u}}_{i+1}, L_{i+1}, \mathbf{z}_i]$ was set before $\operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ and that the verification condition is satisfied, is at most $Q_i Q_{i+1} / |\operatorname{Ch}|$, where Q_i 's are defined as above. By the definition of the verification procedure, it holds that $\tilde{\mathbf{u}}_{i-1} = \tilde{\mathbf{u}}_i - (\bar{\mathbf{A}}\mathbf{z}_i - c_i\mathbf{t}_i) \mod q$. However, because both $\operatorname{HT}[\tilde{\mathbf{u}}_{i-1}, L_{i-1}, \mathbf{z}_{i-2}]$ and $\operatorname{HT}[\tilde{\mathbf{u}}_{i+1}, L_{i+1}, \mathbf{z}_i]$ have been already set before $c_i = \operatorname{HT}[\tilde{\mathbf{u}}_i, L_i, \mathbf{z}_{i-1}]$ is sampled, for fixed $\tilde{\mathbf{u}}_i, \tilde{\mathbf{u}}_{i-1}, \mathbf{z}_i, \mathbf{t}_i, \bar{\mathbf{A}}$, the probability that fresh c_i meets the verification condition is

$$\Pr_{\substack{c_i \overset{\$}\leftarrow \mathsf{Ch}}} \left[\tilde{\mathbf{u}}_{i-1} = \tilde{\mathbf{u}}_i - (\bar{\mathbf{A}}\mathbf{z}_i - c_i\mathbf{t}_i) \bmod q \right],$$

which is at most $1/|\mathsf{Ch}|$. Because there are at most Q_{i+1} different existing entries of $\mathsf{HT}[\tilde{\mathbf{u}}_{i+1}, L_{i+1}, \mathbf{z}_i]$ and fresh c_i is sampled at most Q_i times, by the union bound, we obtain the overall upper bound $Q_iQ_{i+1}/|\mathsf{Ch}|$. The base case i=1 is clearly true: if the tuple $(\tilde{\mathbf{u}}_1, L_1, \mathbf{z}_0)$ is queried after $\mathsf{HT}[\tilde{\mathbf{u}}_2, L_2, \mathbf{z}_1]$ has been set, for the verification condition to be met fresh c_1 must satisfy $c_1\mathbf{t}_1 = \tilde{\mathbf{u}}_1 - \bar{\mathbf{A}}\mathbf{z}_1 \bmod q$ for fixed $\tilde{\mathbf{u}}_1, \mathbf{z}_1, \mathbf{t}_1$. The overall probability is thus bounded by $Q_1Q_2/|\mathsf{Ch}|$ using the same argument as above. All in all, we have that $\Pr[\mathsf{bad}_{\mathsf{ord}}] \leq (1 + \sum_{i=1}^{N-1} (Q_i + Q_i Q_{i+1}))/|\mathsf{Ch}|$. Note that

$$\begin{split} \Pr[\mathsf{bad}_{\mathsf{zcol}}] + \Pr[\mathsf{bad}_{\mathsf{ord}}] &\leq \sum_{i=1}^{N} Q_i(Q_i + 1)/(2|\mathsf{Ch}|) + \left(1 + \sum_{i=1}^{N-1} (Q_i + Q_i Q_{i+1})\right)/|\mathsf{Ch}| \\ &< \left(\sum_{i=1}^{N} Q_i + 1\right)^2/|\mathsf{Ch}| = (Q_h + Q_s + 1)^2/|\mathsf{Ch}|. \end{split}$$

Putting all the bounds above together, we obtain the concrete bound in the theorem statement.

Table 1: Tipping point N_0 (where aggregate signatures start to be smaller than trivial concatenations) and some τ values of our FSwA-SAS (Alg. 2) for the three different parameter sets of Dilithium.

Parameter	Level 2	Level 3	Level 5
\overline{q}	8380417	8380417	8380417
n	256	256	256
(k,ℓ)	(4, 4)	(6, 5)	(8,7)
$B = \gamma - \kappa \cdot \eta$	130994	524092	524168
$\overline{N_0}$	92	138	184
$\tau(200)$	0.9961	0.9985	0.9997
$\tau(250)$	0.9954	0.9979	0.9991
$\tau(500)$	0.9941	0.9966	0.9978
$\tau(1000)$	0.9934	0.9959	0.9971
$\tau(1,000,000)$	0.9927	0.9952	0.9965

Performance Estimates and Comparison

Performance Estimates 4.1

In the following, we provide some concrete sample parameters and performance estimates for the FSwA-SAS from Section 3. We provide a formula for the compression rate τ and a lower bound for N, from which on our SAS signature σ_N is smaller than the trivial solution of concatenating N independent single signatures σ_{con} . The compression rate is defined as $\tau(N) = \frac{\mathsf{len}(\sigma_N)}{\mathsf{len}(\sigma_{\mathsf{con}})}$, where $\mathsf{len}(\cdot)$ denotes the bit size of an element.

A FSwA-SAS signature after N steps is given by $\sigma_N = (\tilde{\mathbf{u}}_N, \mathbf{z}_1, \dots, \mathbf{z}_N)$ and the concatenation of N single FSwA-S signatures by $\sigma_{con} = (c_1, \dots, c_N, \mathbf{z}_1, \dots, \mathbf{z}_N)$. Here, we have applied the standard trick to shorten FSwA-S signatures by replacing the commitment \mathbf{u} by the challenge c. Thus, its compression rate is

$$\tau(N) = \frac{\operatorname{len}(\mathbf{u}) + N \cdot \operatorname{len}(\mathbf{z})}{N \cdot \operatorname{len}(c) + N \cdot \operatorname{len}(\mathbf{z})} = \frac{kn\lceil \log_2 q \rceil + N(k+\ell)n\lceil \log_2 B \rceil}{Nn + N(k+\ell)n\lceil \log_2 B \rceil}$$

$$= 1 - \frac{1}{1 + (k+\ell)\lceil \log_2 B \rceil} + \frac{k\lceil \log_2 q \rceil}{N + N(k+\ell)\lceil \log_2 B \rceil},$$

$$(4)$$

$$=1-\frac{1}{1+(k+\ell)\lceil\log_2 B\rceil}+\frac{k\lceil\log_2 q\rceil}{N+N(k+\ell)\lceil\log_2 B\rceil},$$
(4)

where $\mathbf{u} \in R_q^k$, $\mathbf{z} \in S_B^{\ell+k}$ and $c \in \mathsf{Ch} := \{c \in R : \|c\|_{\infty} = 1 \land \|c\|_1 = \kappa\}$. An element $c \in \mathsf{Ch}$ can be represented by n bits [LDK⁺20, Sec. 5.3].

The SAS signature starts to be smaller than the concatenation as soon as $len(\mathbf{u}) < N \cdot len(c)$, hence, the tipping point is $N_0 > \frac{kn\lceil \log_2 q \rceil}{n} = k\lceil \log_2 q \rceil$. In Table 1, we provide concrete numbers for N_0 and τ for different parameter sets. More

precisely, we take the same parameters as the ones provided for different security levels of Dilithium, denoted by Level 2, Level 3 and Level 5. We can clearly see that in Equation 4 the compression rate asymptotically goes towards $1 - 1/(1 + (k + \ell)\lceil \log_2 B \rceil)$ and for example for the Level 2 parameters of Dilithium this is exactly the rate 0.9927 that we observe at N=1,000,000.

Unlike other proposals to aggregate lattice-based signatures (either interactive [DOTT21, BTT22] or non-interactive [BR21]), the modulus q doesn't need to be increased in our construction. This is due to the fact that we aggregate over the **u**-parts of the signature (which are uniform modulo q), and not over the **z**-parts (which are small and hence the size of their sum increases).

We remark that the needed time to sequentially aggregate N signatures is linear in N. This is unavoidable when sequentially aggregating, as signing cannot be parallelized. As mentioned before, our sequential aggregate signature scheme FSwA-SAS can be seen as the vanilla version of Dilithium, ignoring several optimizations of the latter to further improve efficiency. Hence, the given numbers in Table 1 are only valid for our scheme and do not directly apply to Dilithium.

4.2Comparison With SAS Using Trapdoors

In this part, we compare our lattice-based SAS scheme with existing proposals of lattice-based SAS schemes [EB14, WW19]. As summarized in the introduction, they can be seen as sequential aggregate versions of GPV-signatures. In the following, we take Falcon as a concrete instantiation for such a signature.

As for the FSwA-S signature, the size of a single GPV-signature can be significantly reduced by applying a small trick. More precisely, a Falcon signature of a message m is defined as $\sigma=(s_1,s_2,r)$, where $(s_1,s_2) \in R \times R$ is a pair of short polynomials such that $s_1h+s_2=\mathsf{H}(m,r)$, where H is a random oracle, r is some randomness salt and $h \in R_q$ defines the public basis of the underlying NTRU lattice. Here, R is again the ring $\mathbb{Z}[X]/\langle X^n+1\rangle$ for n a power of 2 and q some prime integer. As s_2 is determined by m and s_1 (given the public key h and the salt r), one can omit s_2 in the signature and only set $\sigma=(s_1,r)$. Intuitively, this (roughly) halves the signature size.

Unfortunately, this trick can't be used in the (sequential) aggregate signature setting. Thus, when assessing the compactness of an aggregate signature, one has to compare it with the trivial concatenation of all single signatures, where each is only composed of the second polynomial. This fair comparison has been done in [EB14], but not in [WW19].

Recently, Espitau et al. [ETWY22] used exactly this trick to make Falcon signatures even shorter. By using elliptical instead of spherical Gaussians, the norm of s_1 can be made smaller. At the same time, the norm of s_2 gets larger, accordingly. Again, this trick does not apply to (sequential) aggregate signatures, as the total size of (s_1, s_2) stays the same.

In the existing SAS schemes that aggregate GPV-style signatures, the main bottleneck is that in the lattice setting there are no known trapdoor permutations. To circumvent this, they replace the trapdoor permutations from the RSA setting by so-called preimage sampleable trapdoor functions [GPV08]. However, those functions have different domain Do and range Ra spaces. In the case of Falcon, the domain is given by R_q , i.e., any element $x \in \mathsf{Ra}$ is of bit length $\mathsf{len}(x) = n\lceil \log_2 q \rceil$. The range, however, is given by pairs of polynomials of degree less than n with coefficients that come from a discrete Gaussian distribution. Naively, one could apply the Gaussian tail bound to argue that the coefficient's absolute values are bounded by some parameter β , and hence any element $y \in Do$ can be represented by a bit string of length $len(y) = 2n \lceil log_2 \beta \rceil$. The specifications of Falcon [PFH⁺20, Sec. 3.11.2] propose a more intelligent representation of elements in the domain by using the Huffman encoding. Note that in both cases it yields len(Ra) > len(Do). As the output of one preimage sampleable function serves as the input for the next preimage sampleable function (of the same domain as before), existing constructions [EB14] pack as many bits of the so-far signature as they can into a vector that serves as the new input. The remaining bits (b := len(Ra) - len(Do))are stored in some vector α and appended (at every step) to the so-far signature and appear at the end in the final sequential aggregate signature. Clearly, they can't achieve a constant-size aggregate

For concreteness, take the sample parameters of Falcon-512, i.e., q=12289 and n=512. It yields $\operatorname{len}(\mathsf{Do})=n\lceil\log_2 q\rceil=7168$ and $\operatorname{len}(\mathsf{Ra})=2\cdot 5000$, and hence $b=2832.^5$ The final sequential aggregate signature after N steps is given by $\sigma_N=(s_{N,1},s_{N,2},\alpha_1,\ldots,\alpha_{N-1},r_1,\ldots,r_N)$, where $(s_{N,1},s_{N,2})\in\mathsf{Ra}$, $\operatorname{len}(\alpha_i)=b$ and $\operatorname{len}(r_i)=328$ for $i\in[N]$. On the other side, the concatenation of N single Falcon signatures is given by $\sigma_{\mathsf{con}}=(s_{1,1},\ldots,s_{N,1},r_1,\ldots,r_N)$, where we applied the 'omit the second polynomial' trick. Thus, its compression rate is given by

$$\begin{split} \tau(N) &= \frac{\mathsf{len}(s_1) + \mathsf{len}(s_2) + (N-1) \cdot \mathsf{len}(\alpha) + N \cdot \mathsf{len}(r)}{N \cdot \mathsf{len}(s_1) + N \cdot \mathsf{len}(r)} \\ &= \frac{2 \cdot 5000 + (N-1)b + N \cdot 328}{N \cdot 5000 + N \cdot 328} = 1 - \frac{5000 - b}{5000 + 328} + \frac{10000 - b}{N(5000 + 328)}, \end{split}$$

where $(s_1, s_2) \in \mathsf{Ra}$, α is the carry-over information and r the salt.

We provide the number N_0 and some τ values for the two different security parameters of Falcon in Table 2. From the equations above, we can clearly see that the compression rate asymptotically goes towards 1 - (5000 - b)/5328, which is exactly the rate 0.5931 that we observe at N = 1,000,000.

In the following, we explain how the recent results of Espitau et al. [ETWY22] extremely leverage the benefit of GPV-style SAS. Overall, they significantly reduce the size of the trivial concatenation by replacing spherical Gaussians by elliptical Gaussians. The main idea is that there are now two different lengths, $len(s_1)$ and $len(s_2)$, where the first holds for s_1 and the latter for s_2 for every

⁴ This analysis has been done by [EB14].

⁵ We compute len(Ra) as $2 \cdot (8 \cdot \text{sbytelen} - 328)$ with sbytelen taken from [PFH⁺20, Table 3.3].

Table 2: Tipping point N_0 (where aggregate signatures start to be smaller than trivial concatenations) and some τ values for SAS based on Falcon-512, for spherical and elliptical Gaussians (distortion factor $\gamma = 8$).

Parameter	Falcon-512 (spherical	l) Falcon-512 (elliptical)
\overline{q}	12289	12289
n	512	512
len(Do)	7168	7168
$len(s_1)$	5000	2952
$len(s_2)$	5000	7048
b	2832	2832
len(r)	328	328
$\overline{N_0}$	4	60
$\tau(150)$	0.6021	0.9780
$\tau(200)$	0.5998	0.9743
$\tau(250)$	0.5985	0.9722
$\tau(500)$	0.5958	0.9678
$\tau(1000)$	0.5944	0.9656
$\tau(1,000,000)$	0.5931	0.9634

pair $(s_1, s_2) \in \mathsf{Ra}$. Whereas before both s_1 and s_2 followed a Gaussian distribution of width σ , they now introduce a distortion factor γ and set $\sigma_1 = \sigma/\gamma$ and $\sigma_2 = \sigma\gamma$. One can see that the total size of (s_1, s_2) is preserved as it yields $2\log_2\sigma = \log_2\sigma_1 + \log_2\sigma_2$. If one takes $\gamma = 8$ (as suggested by Espitau et al. [ETWY22, Table 1]), one can see that $\mathsf{len}(s_1) = 2952$, by again using the formulas of the Falcon specifications.

The compression rate in the elliptical Gaussian case is

$$\begin{split} \tau(N) &= \frac{\mathsf{len}(s_1) + \mathsf{len}(s_2) + (N-1) \cdot \mathsf{len}(\alpha) + N \cdot \mathsf{len}(r)}{N \cdot \mathsf{len}(s_1) + N \cdot \mathsf{len}(r)} \\ &= \frac{2 \cdot 5000 + (N-1)b + N \cdot 328}{N \cdot 2952 + N \cdot 328} = 1 - \frac{2952 - b}{2952 + 328} + \frac{10000 - b}{N(2952 + 328)}, \end{split}$$

where $(s_1, s_2) \in \mathsf{Ra}$, α is the carry-over information and r the salt. Here, we can clearly see that the compression rate asymptotically goes towards 1 - (2952 - b)/(2952 + 328), which is exactly the rate 0.9634 that we observe at N = 1,000,000.

5 Attacks on Existing Schemes

5.1 Attack on [WW19]

In the following, we identify an insecurity of the history-free sequential aggregate signature from Wang and Wu [WW19], published in the proceedings of the PROVSEC conference from 2019. More precisely, Lemma 5.1 gives an attack that breaks its security in the history-free setting. Intuitively, a history-free SAS does not require each signer i to take a so-far message-key pair list L_{i-1} as input. The winning condition in the HF-UF-CMA game (formally recalled in Appendix A.1) is adjusted such that they win as long as the forged message associated with a challenge public key pk has never been queried to the signing oracle. From a high level perspective, the signing procedure of [WW19] closely follows SAS₂ of [GOR18] in the so-called ideal cipher model, except one crucial optimization that reduces the signature size: it deterministically derives ephemeral randomness from the message to be signed and an aggregate so-far, whereas the original SAS₂ requires each signer to append fresh randomness to an aggregate signature. We observe this small change does not sufficiently randomize the scheme and leads to a variant of simple forgery attacks in the history-free setting, which were already pointed out by Brogle et al. [BGR12, App. A] and Gentry et al. [GOR18, Sec. 4.3]. Recall that their construction focuses on lattice signatures that follow the GPV-paradigm. For simplicity, we adapt in the rest of the section the syntax of Falcon, as in Section 5 of [WW19].

Let us (again) briefly recap how Falcon works. As before, we are working over the ring $R_q = \mathbb{Z}_q[X]/\langle X^n+1\rangle$ for some power-of-two integer n and some prime modulus q. The key generation

Algorithm 4: Description of History-Free SAS' [WW19] The message space is $M = \{0, 1\}^l$. The two random oracles are $H_1, H_2: \{0, 1\}^* \to \{0, 1\}^{\lambda}$ and the ideal cipher is $\pi: \{0,1\}^* \times R_q \to R_q$ with inverse $\pi^{-1}: \{0,1\}^* \times R_q \to R_q$. Let $\sigma_0 = ((0,0),0)$. $Gen(1^{\lambda})$ $\mathsf{SeqVerify}(L_N, \sigma_N)$ 1: $(h, T_h) \leftarrow \mathsf{TrapGen}(1^{\lambda})$ 1: $\{(h_1, m_1), \dots, (h_N, m_N)\} := L_N$ 2: pk := h2: $(x_N, \alpha_N) := \sigma_N$ 3: $(s'_1, ..., s'_N) := \alpha_N$ 4: **for** i = N, ..., 1 **do** 3: $sk := (h, T_h)$ 4: return (pk, sk) $\alpha_i = (s_1', \dots, s_i')$ $\mathsf{SeqSign}(\mathsf{sk}_i, m_i, \sigma_{i-1})$ $K_i = h_i || \mathsf{H}_1(m_i) || \mathsf{H}_2(\alpha_i)$ 1: $(x_{i-1}, \alpha_{i-1}) := \sigma_{i-1}$ $s_{i-1} = \pi(K_i, f_{h_i}(x_i))$ 7: 2: $(h_i, T_{h_i}) := \mathsf{sk}_i$ $x_{i-1} = (s_{i-1}, s'_{i-1})$ 8: 3: $K_i := h_i || \mathsf{H}_1(m_i) || \mathsf{H}_2(\alpha_{i-1})$ 9: if $x_{i-1} \notin Do then$ 4: $(s_{i-1}, s'_{i-1}) := x_{i-1}$ 10: return 05: $\alpha_i := \alpha_{i-1} || s'_{i-1}$ 6: $y_i := \pi^{-1}(K_i, s_{i-1}) \in R_q$ 11: **if** $x_0 = (0,0)$ **then** 12: return 1 7: $x_i \leftarrow \mathsf{SamplePre}(T_{h_i}, y_i) \in \mathsf{Do}$ 13: **else** 8: **return** $\sigma_i := (x_i, \alpha_i)$ return 0

algorithm invokes a function TrapGen which outputs a ring element $h \in R_q^6$, together with an associated trapdoor T_h . This trapdoor is needed to invert the function $f_h \colon \mathsf{Do} \subset R_q \times R_q \to R_q = \mathsf{Ra}$, where $f_h(s,s') = hs + s'$, with the help of a pre-image sampleable function $\mathsf{SamplePre}(T_h,\cdot)$. Without specifying the domain Do precisely, we remark that it only contains pairs of short ring elements. The trapdoor defines the secret key, whereas the element h defines the public key. In order to sign a message m, a random oracle $\mathsf{H} \colon \{0,1\}^* \to R_q$ is invoked on m which outputs a ring element in R_q . Then, the function $\mathsf{SamplePre}$ is used to compute $(s,s') \in \mathsf{Do}$ such that $f_h(s,s') = hs + s' = \mathsf{H}(m)$. The signature is defined as x = (s,s'). In order to verify a signature x = (s,s') for a message m, one simply checks if $(s,s') \in \mathsf{Do}$ and if the equation $hs + s' = \mathsf{H}(m)$ holds in R_q .

The main idea of the history-free sequential aggregate signature SAS' by Wang and Wu [WW19] is to adapt the framework for trapdoor-permutation-based sequential aggregate signatures by Gentry et al. [GOR18] to the lattice setting. As in [GOR18], the scheme is making use of an ideal cipher. Additionally, and in contrast to [GOR18], the scheme in [WW19] also uses two random oracles. As the domain $Do \subset R_q \times R_q$ is larger than the range $Ra = R_q$, we don't have trapdoor permutations in the case of lattice signatures, but only pre-image sampleable functions [GPV08]. This is why a so-far signature σ_{i-1} has to be split into a first part (denoted by x_{i-1}) that contains the output of a previous call on SamplePre, and a second part (denoted by α_{i-1}) which stores the information of the previous signatures that didn't fit into the sequential signing process. This part grows linearly in the number of signed messages.

We summarize the SAS' scheme of [WW19] in Alg. 4 (assuming **enc** and **dec** are instantiated with simple split and merge functions as in [WW19, §5]) ⁷ and present the attack in Lemma 5.1. The key idea of the attack is that an adversary can *predict* the one-time key K_i for a message m_i and public key h_i , even though K_i is randomized due to random oracles.

Lemma 5.1. The history-free SAS' described in Algorithm 4 is not HF-UF-CMA.

Proof. Let \mathcal{A} be a PPT adversary. Their goal is to generate an aggregate signature σ^* for a list L^* claiming that signer i signed message m_i (where the public key h_i of signer i is the challenge public key pk given to \mathcal{A}) without having queried the signing oracle OSeqSign on input m_i . \mathcal{A} proceeds as follows.

- 1 Compute $\sigma_{i-1} = (x_{i-1}, \alpha_{i-1})$ for arbitrary and self-chosen key pairs and messages, defining L_i . Let $(s_{i-1}, s'_{i-1}) := x_{i-1}$.
- **2** Choose some $m_i \neq \widetilde{m_i}$ and let $K_i := h_i ||H_1(m_i)||H_2(\alpha_{i-1})$ and $\widetilde{K_i} := h_i ||H_1(\widetilde{m_i})||H_2(\alpha_{i-1})$.

⁶ The ring element $h \in R_q$ is computationally close to uniform assuming the hardness of decision NTRU.

We slightly modified their scheme by adding Step 9-10 in the SeqVerify algorithm in order to check if the so-far signature x_{i-1} recovered at each step lies indeed in the domain Do. Else, it would be very easy for the adversary to come up with forgeries.

```
3 Compute \widetilde{s_{i-1}} := \pi(\widetilde{K_i}, \pi^{-1}(K_i, s_{i-1})).
```

4 Let $\widetilde{x_{i-1}} := (\widetilde{x_{i-1}}, x'_{i-1})$. Query OSeqSign with input $\widetilde{\sigma_{i-1}} := (\widetilde{x_{i-1}}, \alpha_{i-1})$ and $\widetilde{m_i}$.

The oracle responds with $\widetilde{\alpha_i} = (x_i, \alpha_i)$, such that (1) $\alpha_i = \alpha_{i-1} || s'_{i-1}$ and (2) $\pi(\widetilde{K_i}, f_{h_i}(x_i)) = \widetilde{s_{i-1}}$. The adversary outputs $\sigma^* := \widetilde{\alpha_i}$ and $L^* := L_i \cup \{(h_i, m_i)\}$. Recall from Step 3 that $\pi^{-1}(\widetilde{K_i}, \widetilde{s_{i-1}}) = \pi^{-1}(K_i, s_{i-1})$ while m_i has never been queried to OSeqSign. This is a valid forgery as $\pi(K_i, f_{h_i}(x_i)) = s_{i-1}$.

Remark 5.2. An easy fix against this attack, at the expense of larger sequential aggregate signatures, is to make the key for the ideal cipher unpredictable for the adversary. One possible strategy is to freshly sample a truly random string $r \in \{0,1\}^{\lambda}$ and append it to the public key h and the message m to obtain the ideal cipher key K := h||m||r. Besides, this makes the use of the random oracles H_1 and H_2 superfluous. However, the randomness r has to be carried over throughout the sequential signing process, increasing the size of the final signature by $\lambda \cdot N$ bits, where N is the number of involved signatures. This strategy has already been formalized in the second construction of Gentry et al. [GOR18, Sec. 4.2]. One may think that the scheme instantiated with Falcon can be patched by having each signer reject s_{i-1} with too large norm, since maliciously crafted $\widetilde{s_{i-1}}$ in the above attack is uniform in the range R_q of the ideal cipher π . This ad-hoc countermeasure however does not make the scheme provably secure: since the ideal cipher key $K_i = h_i || \mathsf{H}_1(m_i) || \mathsf{H}_2(\alpha_{i-1})$ can still be predicted by querying the random oracle, the ideal cipher table can be determined before any signing query is made, whereas the abort probability analysis in the simulation of aggregate signing oracle in the proof of [WW19, Theorem 1] crucially requires the input of the ideal cipher to be unpredictable. In fact, Gentry et al. instead suggest using tag-based trapdoor permutations [GOR18, Sec. 4.3 as a provable secure way to maintain deterministic signing. It would be interesting to study whether this approach can be adapted to pre-image sampleable functions.

5.2 Attack on [FH20]

In this section, we describe how to mount a (partial) secret-key recovery attack against the Dilithium-based multi-signature by Fukumitsu and Hasegawa [FH20], published in the proceedings of the PROVSEC conference from 2020. In order to be successful, the adversary only needs one valid and honestly generated signatures. The attack exploits the fact that every party of the multi-signature (including the adversary) obtains the full information of the first signature part \mathbf{u} , which enables them to compute $c \cdot \mathbf{s}_2$, where $(\mathbf{s}_1, \mathbf{s}_2)$ is the secret key corresponding to the provided challenge public key. As it is easier to see the vulnerability in the single signature setting, we first describe an insecure variant of FSwA-S (specified in Algorithm 5) and then show how the vulnerability is carried over to the multi-signature from [FH20]. Note that the attack exploits particulars of the compression technique described below. Hence, if the multi-signature scheme of [FH20] is instantiated with plain FSwA-S instead of Dilithium their overall proof strategy in the quantum random oracle model should still hold.⁸

The main difference between the original (and secure) FSwA-S (Algorithm 1) and the modified (and insecure) version (Algorithm 5) is that we apply a trick due to Bai and Galbraith [BG14], which enables to compress the size of the signature. This technique is also used in Dilithium. The key idea is to compute $\mathbf{u} = \mathbf{A}\mathbf{y}$ (instead of $\mathbf{u} = \mathbf{A}\mathbf{y}_1 + \mathbf{y}_2$ as before) and to only use the high order bits HighBits(\mathbf{u}) of \mathbf{u} to derive the challenge c. The function HighBits takes a module element $\mathbf{r} = (r_1, \dots, r_k)^t \in R^k$, decomposes each r_i into the higher-order part $r_{i,1}$ and lower-order part $r_{i,0}$, respectively, and extracts $(r_{1,1}, \dots, r_{k,1})$ (see [LDK+20] for the formal specification). Subsequently, in the verification algorithm only the high order bits of \mathbf{u} and $\mathbf{A}\mathbf{z} - c\mathbf{t}$ are compared to each other. This modification reduces the dimension of \mathbf{z} from $\ell + k$ to ℓ , where $\mathbf{A} \in R_a^{k \times \ell}$.

In contrast to Dilithium and [BG14], our version contains the full commitment \mathbf{u} , and not the challenge c. Transmitting the challenge instead of the commitment is a well-known technique in the lattice setting to further shrink the size of the signature. As we see in the following, it is also crucial for security once the trick of by Bai and Galbraith [BG14] has been applied.

⁸ We thank the authors of [FH20] for clarifying this point.

```
Algorithm 5: Description of Insecure FSwA-S
The challenge space is \mathsf{Ch} := \left\{ c \in R : \ \|c\|_{\infty} = 1 \ \land \ \|c\|_1 = \kappa \right\} and the message space is M = \{0,1\}^l. The random oracle is \mathsf{H} : \{0,1\}^* \to \mathsf{Ch}. Let \mathscr{D}, B and RejSamp be as in Sec. 2.3.
\mathsf{Setup}(1^{\lambda})
                                                                                                                          Sign(sk, m)
         1: \mathbf{A} \stackrel{\$}{\leftarrow} R_q^{k \times \ell}
2: return A
                                                                                                                                   1: (\mathbf{s}_1, \mathbf{s}_2) := \mathsf{sk}
                                                                                                                                   2: \mathbf{t} := \mathbf{A}\mathbf{s}_1 + \mathbf{s}_2 \mod q
                                                                                                                                   3: z := \bot
Gen(A)
                                                                                                                                   4: while z := \bot do
         1: (\mathbf{s}_1, \mathbf{s}_2) \stackrel{\$}{\leftarrow} S_{\eta}^{\ell} \times S_{\eta}^{k}
2: \mathbf{t} := \mathbf{A}\mathbf{s}_1 + \mathbf{s}_2 \bmod q
                                                                                                                                                 \mathbf{y} \overset{\hspace{0.1em}\mathsf{\scriptscriptstyle\$}}{\leftarrow} \mathscr{D}^{\ell}
                                                                                                                                                  \mathbf{u} := \mathbf{A}\mathbf{y} \bmod q
                                                                                                                                                  c := \mathsf{H}(\mathsf{HighBits}(\mathbf{u}), \mathbf{t}, m)
         3: sk := (s_1, s_2)
         4: pk := t
                                                                                                                                                  \mathbf{z} := c \cdot \mathbf{s}_1 + \mathbf{y}
                                                                                                                                                  z := \mathsf{RejSamp}(z, c \cdot s)
         5: return (sk, pk)
                                                                                                                                 10: \sigma := (\mathbf{u}, \mathbf{z})
Ver(pk, \sigma, m)
                                                                                                                                 11: return \sigma
         1: (\mathbf{u}, \mathbf{z}) := \sigma
         2: \mathbf{t} := pk
         3: c := \mathsf{H}(\mathsf{HighBits}(\mathbf{u}), \mathbf{t}, m)
         4: if \|\mathbf{z}\|_{\infty} \leq B \wedge \mathsf{HighBits}(\mathbf{Az} - c \cdot \mathbf{t}) =
               \mathsf{HighBits}(\mathbf{u})\ \mathbf{then}
                     return 1
         6: else
         7:
                       return 0
```

Lemma 5.3. Let $\mathbf{A} \leftarrow \mathsf{Setup}(1^{\lambda})$ and $(\mathsf{sk}, \mathsf{pk}) \leftarrow \mathsf{Gen}(\mathbf{A})$ with $\mathsf{sk} = (\mathbf{s}_1, \mathbf{s}_2)$ as in Alg. 5. Given $(\mathbf{u}, \mathbf{z}) = \sigma \leftarrow \mathsf{Sign}(\mathsf{sk}, m)$ for the messages $m \in M$ together with the public key $\mathsf{pk} = \mathbf{t}$, a PPT adversary \mathcal{A} can recover \mathbf{s}_2 .

Proof. The adversary re-constructs $c = H(HighBits(\mathbf{u}), \mathbf{t}, m)$ and computes

$$\mathbf{b} = \mathbf{A}\mathbf{z} - \mathbf{u} - c\mathbf{t} = \mathbf{A}\mathbf{y} + c\mathbf{A}\mathbf{s}_1 - \mathbf{A}\mathbf{y} - c\mathbf{A}\mathbf{s}_1 - c\mathbf{s}_2 = -c\mathbf{s}_2.$$

Note that the last equation gives $\mathbf{b} = -c\mathbf{s}_2 \mod q$, but as there is no wrapping around modulo q (as both c and \mathbf{s}_2 are short elements) the equation also holds in R. The adversary now embeds the elements in the field $K = \mathbb{Q}[x]/(x^d+1)$ and uses the fact that every non-zero element is invertible in K. Note that $0 \notin \mathsf{Ch}$. Hence, they multiply the result by c^{-1} (the K-inverse of c) and recover $-\mathbf{s}_2$.

We can now easily move to the multi-signature in [FH20]. We don't re-state the full protocol of their scheme here, but simply refer to Figure 5 in [FH20]. For simplicity, we use our notations in the following and ignore the optimization to reduce the public key size in [FH20, Fig. 5]. The main issue is that every signer broadcasts untruncated \mathbf{u} in the clear during the interactive signing process. We now explain how an adversary can during the common signing procedure recover half of the secret key of the single honest signer.

Lemma 5.4. Let MS = (Setup, Gen, Sign, Ver) be the multi-signature as defined in [FH20, Fig. 5]. Further, let \mathcal{A} be a PPT adversary who is controlling all-but-one parties. We denote by (pk*, sk*) the key pair of the honest signer, where $sk^* = (s_1^*, s_2^*)$. After one successful multi-signature signing process, \mathcal{A} can recover s_2^* .

Proof. Without loss of generality, let party 1 be the honest signer and party 2 to party N be the onest controlled by \mathcal{A} . Let \mathbf{t}_i denote the public key of party $i \in [N]$. During the signing process, every party computes $\mathbf{u}_i := \mathbf{A}\mathbf{y}_i$ (denoted by \mathbf{w}_v in the original protocol). At the second stage, they broadcasts \mathbf{u}_i to all the co-signers. In particular, \mathcal{A} receives \mathbf{u}_1 and computes $\mathbf{u} := \sum_{i \in [N]} \mathbf{u}_i$ (denoted by \mathbf{w} in the original protocol). The multi-signature contains $\mathbf{z} := \sum_{i \in [N]} \mathbf{z}_i$, where $\mathbf{z}_i = \mathbf{y}_i + c_i \mathbf{s}_{i,1}$. The adversary re-constructs all challenges $c_i = \mathsf{H}(\mathsf{HighBits}(\mathbf{u}), \mathbf{t}_i, m)$ and computes

$$\mathbf{b} = \mathbf{A}\mathbf{z} - \mathbf{u} - \sum_{i \in [N]} c_i \mathbf{t}_i = \sum_{i \in [N]} \mathbf{A}\mathbf{y}_i + c_i \mathbf{A}\mathbf{s}_{i,1} - \mathbf{A}\mathbf{y}_i - c_i \mathbf{A}\mathbf{s}_{i,1} - c_i \mathbf{s}_{i,2} = -\sum_{i \in [N]} c_i \mathbf{s}_{i,2}.$$

As they know $\mathbf{s}_{i,2}$ for all $1 < i \le N$, they can compute $\mathbf{b} + \sum_{i=2}^{N} c_i \mathbf{s}_{i,2}$ and recover $c_1 \mathbf{s}_{1,2}$. With the same reasoning as in the attack above, they can easily recover $\mathbf{s}_{1,2} = \mathbf{s}_2^*$.

Remark 5.5. In the simple signature setting (Algorithm 5) it is easy to fix this attack by only outputting $\sigma := (\mathsf{HighBits}(\mathbf{u}), \mathbf{z})$. However, in the multi-signature setting, this fix doesn't seem to apply in a trivial manner. The problem is that the function $\mathsf{HighBits}$ is not linear, and in general $\mathsf{HighBits}(\mathbf{u}_1) + \mathsf{HighBits}(\mathbf{u}_2) \neq \mathsf{HighBits}(\mathbf{u}_1 + \mathbf{u}_2)$.

6 Conclusion

In this paper, we proposed a sequential aggregate signature based on the FSwA framework and showed that it can indeed save bandwidth compared to the naive concatenation. It exploits the fact that aggregation of the u-part is much more critical than the z-part in the lattice-based FSwA setting. Admittedly, the benefit of our construction appears still limited in practice according to the concrete parameter analysis of Section 4. It would make an interesting future direction to design a new aggregation paradigm tailored to FSwA exploiting small c_i 's in the non-aggregated part.

Another natural follow-up question would be how to adapt our scheme to half-aggregate Dilithium signatures incorporating all the bit-truncation optimizations. In terms of correctness, we observe no obvious issue because our scheme can be securely modified by aggregating the u-part with XOR instead of summation $\operatorname{mod} q$ (as already observed by [CZ22] in the Schnorr setting), and thus the fact that HighBits destroys homomorphism is not a major obstacle, unlike the issue with [FH20] we pointed out in Section 5.2. However, one must carefully adapt the probability that bad events happen in the reduction, and thus we leave detailed analysis for future work.

We also observe that the situation with Falcon-based SAS is not satisfactory either given that the recent trick of [ETWY22] significantly saves the size of naive concatenation and one of the existing instantiations turned out to be insecure. Since aggregation of hash-then-signatures essentially amounts to batch-proving the knowledge of short preimages of the function $f_h: (s, s') \mapsto hs + s'$ and requires no proof of correct hash evaluation, we conjecture that generic methods proposed by [DGKV22, ACL⁺22] will likely lead to concretely efficient instantiations in this setting.

Acknowledgment

Katharina Boudgoust is supported by the Danish Independent Research Council under project number 0165-00107B (C3PO) as well as by the Protocol Labs Research Grant Program RFP-013. Akira Takahashi is supported by the Protocol Labs Research Grant Program PL-RGP1-2021-064. The authors are grateful for Claudio Orlandi for discussions in the earlier stages of this work. We thank our anonymous referees for their thorough proof reading and constructive feedback.

References

- ACL⁺22. M. R. Albrecht, V. Cini, R. W. F. Lai, G. Malavolta, and S. A. Thyagarajan. Lattice-based SNARKs: Publicly verifiable, preprocessing, and recursively composable. Cryptology ePrint Archive, Report 2022/941, 2022. https://eprint.iacr.org/2022/941. 3, 4, 21
- AGH10. J. H. Ahn, M. Green, and S. Hohenberger. Synchronized aggregate signatures: new definitions, constructions and applications. In ACM CCS 2010, pp. 473–484. ACM Press, 2010. 5
- BG14. S. Bai and S. D. Galbraith. An improved compression technique for signatures based on learning with errors. In CT-RSA 2014, vol. 8366 of LNCS, pp. 28–47. Springer, Heidelberg, 2014. 5, 19
- BGLS03. D. Boneh, C. Gentry, B. Lynn, and H. Shacham. Aggregate and verifiably encrypted signatures from bilinear maps. In EUROCRYPT 2003, vol. 2656 of LNCS, pp. 416–432. Springer, Heidelberg, 2003. 3
- BGOY07. A. Boldyreva, C. Gentry, A. O'Neill, and D. H. Yum. Ordered multisignatures and identity-based sequential aggregate signatures, with applications to secure routing. In ACM CCS 2007, pp. 276–285. ACM Press, 2007. 3
- BGR12. K. Brogle, S. Goldberg, and L. Reyzin. Sequential aggregate signatures with lazy verification from trapdoor permutations - (extended abstract). In ASIACRYPT 2012, vol. 7658 of LNCS, pp. 644–662. Springer, Heidelberg, 2012. 3, 17, 24

- BJRW20. K. Boudgoust, C. Jeudy, A. Roux-Langlois, and W. Wen. Towards classical hardness of module-LWE: The linear rank case. In ASIACRYPT 2020, Part II, vol. 12492 of LNCS, pp. 289–317. Springer, Heidelberg, 2020. 6
- BJRW23. K. Boudgoust, C. Jeudy, A. Roux-Langlois, and W. Wen. On the hardness of module learning with errors with short distributions. *J. Cryptol.*, 36(1):1, 2023. 6
- BK20. D. Boneh and S. Kim. One-time and interactive aggregate signatures from lattices. *preprint*, 2020. 4
- BLR⁺18. S. Bai, T. Lepoint, A. Roux-Langlois, A. Sakzad, D. Stehlé, and R. Steinfeld. Improved security proofs in lattice-based cryptography: Using the Rényi divergence rather than the statistical distance. *Journal of Cryptology*, 31(2):610–640, 2018. 5
- BN06. M. Bellare and G. Neven. Multi-signatures in the plain public-key model and a general forking lemma. In ACM CCS 2006, pp. 390–399. ACM Press, 2006. 3, 4, 5
- BNN07. M. Bellare, C. Namprempre, and G. Neven. Unrestricted aggregate signatures. In *ICALP 2007*, vol. 4596 of *LNCS*, pp. 411–422. Springer, Heidelberg, 2007. 3
- Bol
03. A. Boldyreva. Threshold signatures, multisignatures and blind signatures based on the gap
Diffie-Hellman-group signature scheme. In PKC~2003, vol. 2567 of LNCS, pp. 31–46. Springer,
Heidelberg, 2003. 5
- BR21. K. Boudgoust and A. Roux-Langlois. Compressed linear aggregate signatures based on module lattices. Cryptology ePrint Archive, Report 2021/263, 2021. https://eprint.iacr.org/2021/263. 4, 15
- BTT22. C. Boschini, A. Takahashi, and M. Tibouchi. Musig-l: Lattice-based multi-signature with single-round online phase. *IACR Cryptol. ePrint Arch.*, p. 1036, 2022. Accepted at Crypto 22. 5, 15
- CGKN21. K. Chalkias, F. Garillot, Y. Kondi, and V. Nikolaenko. Non-interactive half-aggregation of EdDSA and variants of Schnorr signatures. In CT-RSA 2021, vol. 12704 of LNCS, pp. 577–608. Springer, Heidelberg, 2021. 3, 4
- Cra96. R. Cramer. Modular Design of Secure yet Practical Cryptographic Protocols. PhD thesis, CWI, Amsterdam, 1996. https://ir.cwi.nl/pub/21438. 3
- CZ22. Y. Chen and Y. Zhao. Half-aggregation of schnorr signatures with tight reductions. In ESORICS 2022, Part II, vol. 13555 of LNCS, pp. 385–404. Springer, Heidelberg, 2022. 3, 4, 8, 10, 21, 24, 25
- DEF⁺19. M. Drijvers, K. Edalatnejad, B. Ford, E. Kiltz, J. Loss, G. Neven, and I. Stepanovs. On the security of two-round multi-signatures. In 2019 IEEE Symposium on Security and Privacy, pp. 1084–1101. IEEE Computer Society Press, 2019. 3
- DGKV22. L. Devadas, R. Goyal, Y. Kalai, and V. Vaikuntanathan. Rate-1 non-interactive arguments for batch-np and applications. Cryptology ePrint Archive, Paper 2022/1236, 2022. https://eprint.iacr.org/2022/1236. 3, 4, 21
- DHSS20. Y. Doröz, J. Hoffstein, J. H. Silverman, and B. Sunar. MMSAT: A scheme for multimessage multiuser signature aggregation. Cryptology ePrint Archive, Report 2020/520, 2020. https://eprint.iacr.org/2020/520. 4
- DOTT21. I. Damgård, C. Orlandi, A. Takahashi, and M. Tibouchi. Two-round n-out-of-n and multi-signatures and trapdoor commitment from lattices. In PKC 2021, Part I, vol. 12710 of LNCS, pp. 99–130. Springer, Heidelberg, 2021. 5, 15
- EB14. R. El Bansarkhani and J. Buchmann. Towards lattice based aggregate signatures. In AFRICACRYPT 14, vol. 8469 of LNCS, pp. 336–355. Springer, Heidelberg, 2014. 4, 15, 16
- ETWY22. T. Espitau, M. Tibouchi, A. Wallet, and Y. Yu. Shorter hash-and-sign lattice-based signatures. *IACR Cryptol. ePrint Arch.*, p. 785, 2022. Accepted at Crypto 22. 5, 16, 17, 21
- FH20. M. Fukumitsu and S. Hasegawa. A lattice-based provably secure multisignature scheme in quantum random oracle model. In *ProvSec 2020*, vol. 12505 of *LNCS*, pp. 45–64. Springer, Heidelberg, 2020. 2, 5, 19, 20, 21
- FLS12. M. Fischlin, A. Lehmann, and D. Schröder. History-free sequential aggregate signatures. In SCN 12, vol. 7485 of LNCS, pp. 113–130. Springer, Heidelberg, 2012. 3
- FS87. A. Fiat and A. Shamir. How to prove yourself: Practical solutions to identification and signature problems. In *CRYPTO'86*, vol. 263 of *LNCS*, pp. 186–194. Springer, Heidelberg, 1987. 3
- FSZ22. N. Fleischhacker, M. Simkin, and Z. Zhang. Squirrel: Efficient synchronized multi-signatures from lattices. *IACR Cryptol. ePrint Arch.*, p. 694, 2022. 5
- GLP12. T. Güneysu, V. Lyubashevsky, and T. Pöppelmann. Practical lattice-based cryptography: A signature scheme for embedded systems. In CHES 2012, vol. 7428 of LNCS, pp. 530–547. Springer, Heidelberg, 2012. 6, 7

- GOR18. C. Gentry, A. O'Neill, and L. Reyzin. A unified framework for trapdoor-permutation-based sequential aggregate signatures. In *PKC 2018*, *Part II*, vol. 10770 of *LNCS*, pp. 34–57. Springer, Heidelberg, 2018. 3, 7, 17, 18, 19, 24
- GPV08. C. Gentry, C. Peikert, and V. Vaikuntanathan. Trapdoors for hard lattices and new cryptographic constructions. In 40th ACM STOC, pp. 197–206. ACM Press, 2008. 4, 16, 18
- GR06. C. Gentry and Z. Ramzan. Identity-based aggregate signatures. In PKC 2006, vol. 3958 of LNCS, pp. 257–273. Springer, Heidelberg, 2006. 5
- HKW15. S. Hohenberger, V. Koppula, and B. Waters. Universal signature aggregators. In *EURO-CRYPT 2015, Part II*, vol. 9057 of *LNCS*, pp. 3–34. Springer, Heidelberg, 2015. 3
- Kas22. Y. Kondi and abhi shelat. Improved straight-line extraction in the random oracle model with applications to signature aggregation. Cryptology ePrint Archive, Paper 2022/393, 2022. https://eprint.iacr.org/2022/393. 3
- LDK⁺20. V. Lyubashevsky, L. Ducas, E. Kiltz, T. Lepoint, P. Schwabe, G. Seiler, D. Stehlé, and S. Bai. CRYSTALS-DILITHIUM. Technical report, National Institute of Standards and Technology, 2020. available at https://csrc.nist.gov/projects/post-quantum-cryptography/post-quantum-cryptography-standardization/round-3-submissions. 4, 5, 15, 19
- LMRS04. A. Lysyanskaya, S. Micali, L. Reyzin, and H. Shacham. Sequential aggregate signatures from trapdoor permutations. In EUROCRYPT 2004, vol. 3027 of LNCS, pp. 74–90. Springer, Heidelberg, 2004. 3, 7
- LOS⁺06. S. Lu, R. Ostrovsky, A. Sahai, H. Shacham, and B. Waters. Sequential aggregate signatures and multisignatures without random oracles. In *EUROCRYPT 2006*, vol. 4004 of *LNCS*, pp. 465–485. Springer, Heidelberg, 2006. 3
- LPR13. V. Lyubashevsky, C. Peikert, and O. Regev. A toolkit for ring-LWE cryptography. In *EURO-CRYPT 2013*, vol. 7881 of *LNCS*, pp. 35–54. Springer, Heidelberg, 2013. 6, 9
- LS15. A. Langlois and D. Stehlé. Worst-case to average-case reductions for module lattices. *Des. Codes Cryptogr.*, 75(3):565–599, 2015. 6, 7
- Lyu
09. V. Lyubashevsky. Fiat-Shamir with aborts: Applications to lattice and factoring-based signatures.
 In ASIACRYPT 2009, vol. 5912 of LNCS, pp. 598–616. Springer, Heidelberg, 2009. 4, 6
- Lyu12. V. Lyubashevsky. Lattice signatures without trapdoors. In EUROCRYPT 2012, vol. 7237 of LNCS, pp. 738–755. Springer, Heidelberg, 2012. 4, 6, 7
- Mic18. D. Micciancio. On the hardness of learning with errors with binary secrets. Theory Comput., $14(1):1-17,\ 2018.\ 9$
- MOR01. S. Micali, K. Ohta, and L. Reyzin. Accountable-subgroup multisignatures: Extended abstract. In *ACM CCS 2001*, pp. 245–254. ACM Press, 2001. 5
- Nev
08. G. Neven. Efficient sequential aggregate signed data. In EUROCRYPT~2008, vol. 4965 of LNCS, pp. 52–69. Springer, Heidelberg, 2008. 3
- NRS21. J. Nick, T. Ruffing, and Y. Seurin. MuSig2: Simple two-round Schnorr multi-signatures. In CRYPTO 2021, Part I, vol. 12825 of LNCS, pp. 189–221, Virtual Event, 2021. Springer, Heidelberg. 3
- PFH⁺20. T. Prest, P.-A. Fouque, J. Hoffstein, P. Kirchner, V. Lyubashevsky, T. Pornin, T. Ricosset, G. Seiler, W. Whyte, and Z. Zhang. FALCON. Technical report, National Institute of Standards and Technology, 2020. available at https://csrc.nist.gov/projects/post-quantum-cryptography/post-quantum-cryptography-standardization/round-3-submissions. 4, 16
- Sch91. C.-P. Schnorr. Efficient signature generation by smart cards. *Journal of Cryptology*, 4(3):161–174, 1991. 3
- vEH14. T. van Erven and P. Harremoës. Rényi divergence and kullback-leibler divergence. *IEEE Trans.* Inf. Theory, 60(7):3797–3820, 2014. 5
- WW19. Z. Wang and Q. Wu. A practical lattice-based sequential aggregate signature. In *ProvSec 2019*, vol. 11821 of *LNCS*, pp. 94–109. Springer, Heidelberg, 2019. 2, 4, 5, 15, 16, 17, 18, 19
- WW22. B. Waters and D. J. Wu. Batch arguments for sfNP and more from standard bilinear group assumptions. In CRYPTO 2022, Part II, vol. 13508 of LNCS, pp. 433–463. Springer, Heidelberg, 2022. 3

```
Game 2: Description of the HF-UF-CMA<sub>SAS'</sub> (A, \lambda) Security Game
1: pp \leftarrow Setup(1^{\lambda})
                                                                                  \mathsf{OSeqSign}(m_i, \sigma_{i-1})
2: (pk, sk) \leftarrow Gen(pp)
                                                                                         1: \sigma_i \leftarrow \mathsf{SeqSign}(\mathsf{sk}, m_i, \sigma_{i-1})
3: \mathcal{Q} := \emptyset
                                                                                        2: Q := Q \cup \{m_i\}
4: (L_N^*, \sigma_N^*) \leftarrow \mathcal{A}^{\mathsf{OSeqSign}}(\mathsf{pp}, \mathsf{pk})
                                                                                        3: return \sigma_i
5: if SeqVerify(L_N^*,\sigma_N^*) \wedge \exists i^* \in [N]: (\mathsf{pk}_{i^*} =
     pk \wedge m_{i^*} \notin Q) then
           return 1
6:
7: else
           return 0
8:
```

A More Security Notions for Sequential Aggregate Signatures

In this section we discuss alternative flavors of security notions for SAS. Gentry et al. [GOR18] distinguish between the *full-history* (see Definition 2.9) and the *history-free* case. Additionally, we formalize yet another flavor as described by Chen and Zhao [CZ22], which we call *partial-signature history-free*.

A.1 Histrory-Free Sequential Aggregate Signatures

History-free sequential aggregate signatures (SAS') were first introduced by Brogle et al. [BGR12]. We recall now their syntax, together with the definition of security following Gentry et al. [GOR18]. Note that the only difference to the full-history setting is that now the sequential signing algorithm SeqSign doesn't take as input the list L_{i-1} of public keys and messages (i.e., the 'history') which have been used to compute the so-far signature σ_{i-1} . Consequently, the winning condition in the security game changes accordingly. In the full-history case, a valid forgery could use an already queried message, as long as the history changed. Now, a valid forgery must be on a non-queried message.

Definition A.1 (SAS'). A history-free sequential aggregate signature (SAS') for a message space M consists of a tuple of PPT algorithms SAS' = (Setup, Gen, SeqSign, SeqVerify) defined as follows:

Setup, Gen and SeqVerify as in Definition 2.7.

 $\mathsf{SeqSign}(\mathsf{sk}_i, m_i, \sigma_{i-1}) \to \sigma_i$: On input a secret key sk_i , a message $m_i \in M$, and a so-far signature σ_{i-1} , the sequential signing algorithm outputs a new so-far signature σ_i .

Definition A.2 (HF-UF-CMA Security). A SAS' scheme satisfies <u>h</u>istory-<u>f</u>ree <u>unf</u>orgeabilty against <u>c</u>hosen <u>m</u>essage <u>a</u>ttacks, if for all PPT adversaries A,

$$\mathsf{Adv}_{\mathsf{SAS'}}^{\mathsf{HF-UF-CMA}}(\mathcal{A}) := \Pr \left[\mathsf{HF-UF-CMA}_{\mathsf{SAS'}}(\mathcal{A}, \lambda) = 1 \right] = \mathsf{negl}(\lambda),$$

where the HF-UF-CMA_{SAS'} game is described in Game 2.

A.2 Partial-Signature History-Free Sequential Aggregate Signatures

As already mentioned before, our construction of Section 3 is inspired by the paper on sequential half-aggregation of Schnorr signatures by Chen and Zhao [CZ22]. In their work, they proposed yet another security model, without providing a formal definition (and name) for it. Their model is specifically tailored to Schnorr type signatures. In the following, we formalize the security model by trying to be as general as possible and hence making the notion useful also for other type of signatures.

Recall that in both the history-free and the full-history case (Def. 2.9 and Def. A.2), the sequential aggregate signature takes the full description of the so-far signature σ_{i-1} as input to the SeqSign algorithm and then outputs the full description of the next so-far signature σ_i .

We now define a variant of the history-free case (i.e., no list L_{i-1} given to SeqSign) where the amount of information that needs to be send to the next signer is reduced, at the expense of introducing a new role, that we call the Combine algorithm. At each signing step, the signer forwards only a partial description of σ_{i-1} to the next signer, while at the same time sending some complementary description of σ_{i-1} to the combiner. Note that the complementary information can overlap with the partial description. At the end, the combiner takes the partial description of the final signature and all the complementary information they have received so far, to derive the full description of the final signature.

We call it the *partial-signature history-free* case. The winning condition for the adversary stays the same as in the history-free case.

Chen and Zhao [CZ22] actually consider a stronger notion where the winning condition is relaxed to allow for forgeries that are on already queried messages as long as the corresponding complementary part (in their case the response of the underlying Σ -protocol) is new. We suggest calling this the *strong* PS-HF-UF-CMA security.

Definition A.3 (SAS"). A partial-signature history-free sequential aggregate signature scheme (SAS") for a message space M consists of a tuple of PPT algorithms SAS" = (Setup, Gen, SeqSign, SeqVerify, Combine) defined as follows:

Setup, Gen and SeqVerify as in Definition 2.7 and A.1.

SeqSign($\operatorname{sk}_i, m_i, \operatorname{Part}(\sigma_{i-1})$) $\to (\operatorname{Part}(\sigma_i), \operatorname{Compl}(\sigma_i))$: On input a secret key sk_i, a message $m_i \in M$, and a partial description $\operatorname{Part}(\sigma_{i-1})$ of the so-far signature σ_{i-1} , the sequential signing algorithm outputs a partial description $\operatorname{Part}(\sigma_i)$ of a new so-far signature σ_i and some complementary information $\operatorname{Compl}(\sigma_i)$.

 $\mathsf{Combine}(\mathsf{Compl}(\sigma_1), \dots, \mathsf{Compl}(\sigma_{N-1}), \mathsf{Part}(\sigma_N)) \to \sigma_N.$

Definition A.4 (PS-HF-UF-CMA Security). A SAS" scheme satisfies \underline{p} artial- \underline{s} ignature \underline{h} istory-free \underline{u} mforgeabilty against \underline{c} hosen \underline{m} essage \underline{a} ttacks, if for all PPT adversaries \mathcal{A} ,

$$\mathsf{Adv}^{\mathsf{PS-HF-UF-CMA}}_{\mathsf{SAS''}}(\mathcal{A}) := \Pr\left[\mathsf{PS-HF-UF-CMA}_{\mathsf{SAS''}}(\mathcal{A},\lambda) = 1\right] = \mathsf{negl}(\lambda),$$

where the PS-HF-UF-CMA_{SAS''} game is described in Game 3. It satisfies strong PS-HF-UF-CMA security if the modifications in purple apply.

The interest of this model lies in the fact that when only half-aggregation is possible (as it is the case in our construction and the one of [CZ22]), then a lot of information in every so-far signature is redundant and simply carried over to the next signing step. In our case, this corresponds to the list of responses $\mathbf{z}_1, \ldots, \mathbf{z}_i$ that is attached to every so-far signature σ_i . One possible instantiation of Part and Compl would be to set $\mathsf{Part}(\sigma_i) = (\tilde{\mathbf{u}}_i, \mathbf{z}_i)$ and $\mathsf{Compl}(\sigma_i) = \mathbf{z}_i$. By outsourcing this redundant information to the combiner, one safes significantly in bandwidth.

Algorithm 6: Description of the FSwA-SAS" Sequential Aggregate Signature

```
The challenge space is \mathsf{Ch} := \left\{ c \in R : \|c\|_{\infty} = 1 \land \|c\|_{1} = \kappa \right\} and the message space is M' = \{0,1\}^{l}.
     The random oracle is H: \{0,1\}^* \to Ch. The starting point is i=1. Let L_0 = \emptyset and \sigma_0 = (\mathbf{0},\mathbf{0}).
      Setup and Gen are as in Algorithm 2.
 \mathsf{SeqSign}(\mathsf{sk}_i, m_i, \mathsf{Part}(\sigma_{i-1}))
                                                                                                          SeqVerify(L_N, \sigma_N)
        1: (\tilde{\mathbf{u}}_{i-1}, \mathbf{z}_{i-1}) := \mathsf{Part}(\sigma_{i-1})
                                                                                                                 1: (\mathbf{t}_1, m_1) || \dots || (\mathbf{t}_N, m_N) := L_N
                                                                                                                 2: (\tilde{\mathbf{u}}_N, \mathbf{z}_1, \dots, \mathbf{z}_N) := \sigma_N
        2: \mathbf{s}_i := \mathsf{sk}_i
       3: \mathbf{t}_i := \mathbf{A}\mathbf{s}_i \bmod q
        4: \mathbf{z}_i := \bot
                                                                                                                 4: if \exists i such that \mathbf{t}_i has no invertible element
        5: while \mathbf{z}_i := \bot \mathbf{do}
                                                                                                                       then
                     \mathbf{y}_i \leftarrow \mathscr{D}^{\ell+k}
                                                                                                                             return 0
                      \mathbf{u}_i := \bar{\mathbf{A}}\mathbf{y}_i \bmod q
                                                                                                                 6: for i = N, ..., 1 do
                      \tilde{\mathbf{u}}_i := \tilde{\mathbf{u}}_{i-1} + \mathbf{u}_i \bmod q
                                                                                                                             if \|\mathbf{z}_i\|_2 > B then
                      c_i := \mathsf{H}(\tilde{\mathbf{u}}_i, \mathbf{t}_i, m_i, \mathbf{z}_{i-1}, i)
                                                                                                                                    return 0
       9:
                                                                                                                 8:
                      \mathbf{z}_i := c_i \cdot \mathbf{s}_i + \mathbf{y}_i
                                                                                                                               c_i := \mathsf{H}(\tilde{\mathbf{u}}_i, \mathbf{t}_i, m_i, \mathbf{z}_{i-1}, i)
                                                                                                                 9:
                                                                                                                               \mathbf{u}_i := \bar{\mathbf{A}}\mathbf{z}_i - c_i\mathbf{t}_i \bmod q
                      \mathbf{z}_i := \mathsf{RejSamp}(\mathbf{z}_i, c_i \cdot \mathbf{s}_i)
                                                                                                                               \tilde{\mathbf{u}}_{i-1} := \tilde{\mathbf{u}}_i - \mathbf{u}_i \bmod q
      12: \mathsf{Part}(\sigma_i) := (\tilde{\mathbf{u}}_i, \mathbf{z}_i)
     13: \mathsf{Compl}(\sigma_i) := \mathbf{z}_i
                                                                                                               12: if \tilde{\mathbf{u}}_1 = \mathbf{u}_1 then return 1
     14: return Part(\sigma_i), Compl(\sigma_i)
 \mathsf{Combine}(\mathsf{Compl}(\sigma_1), \dots, \mathsf{Compl}(\sigma_{N-1}), \mathsf{Part}(\sigma_N))
        1: \mathbf{z}_i := \mathsf{Compl}(\sigma_i) \text{ for } i \in [N-1]
       2: (\tilde{\mathbf{u}}_N, \mathbf{z}_N) := \mathsf{Part}(\sigma_N)
       3: \sigma_N = (\tilde{\mathbf{u}}_N, \mathbf{z}_1, \dots, \mathbf{z}_N)
        4: return \sigma_N
```

Algorithm 6 specifies how our FSwA-SAS from Section 3 can easily be modified to FSwA-SAS" which fulfills PS-HF-UF-CMA security. The main modifications are that in SeqSign and SeqVerify the random oracle H now doesn't take the full history L_i as input, but only the current public key-message pair (\mathbf{t}_i, m_i) together with an index i that fixes the position in the sequence. The combine algorithm simply puts together the relevant information in order to define σ_N .

It is straightforward to adapt the security reduction (as specified in Alg. 3) to the PS-HF-UF-CMA setting. Again, one only has to index of the different tables ZT, HT, MT by replacing the full history L_i by the 'current' history (\mathbf{t}_i, m_i, i) . This change does not affect the bad flag analysis conducted in the proof for Theorem 3.4: $\mathsf{bad}_{\mathsf{mcol}}$ only depends on the size of the message space, $\mathsf{bad}_{\mathsf{ucol}}$ only depends on min-entropy of freshly sampled \mathbf{u}_i , $\mathsf{bad}_{\mathsf{zcol}}$ only depends on the distribution of $\mathbf{u}_i + c_i \mathbf{t}_i$, and $\mathsf{bad}_{\mathsf{ord}}$ only depends on the distribution of $\tilde{\mathbf{u}}_{i-1} - (\bar{\mathbf{A}}\mathbf{z}_i - c_i \mathbf{t}_i)$, respectively. Thus, all bad flags are preserved and occur with the same probability.

Theorem A.5 (PS-HF-UF-CMA security). If the signature scheme FSwA-S with message space $M = \{0,1\}^l$, as described in Algorithm 1, is UF-CMA secure, then is the sequential aggregate signature FSwA-SAS", as described in Algorithm 6, PS-HF-UF-CMA secure.